



OPEN DATA SCIENCE CONFERENCE

Boston | April 30 - May 4th, 2019

@newfront | @odsc



#ODSC

BOSTON
APR 30 - MAY 3

Real-ish Time Predictive Analytics with Spark Structured Streaming

Scott Haines

Principal Software Engineer
Twilio





Introductions

Scott Haines, Principal Software Engineer,
Voice Insights, **Twilio**

<http://bit.ly/odsceastrealtime>

<http://bit.ly/odsc-east-2019-realish-predict>

@newfront | @odsc | @twilio





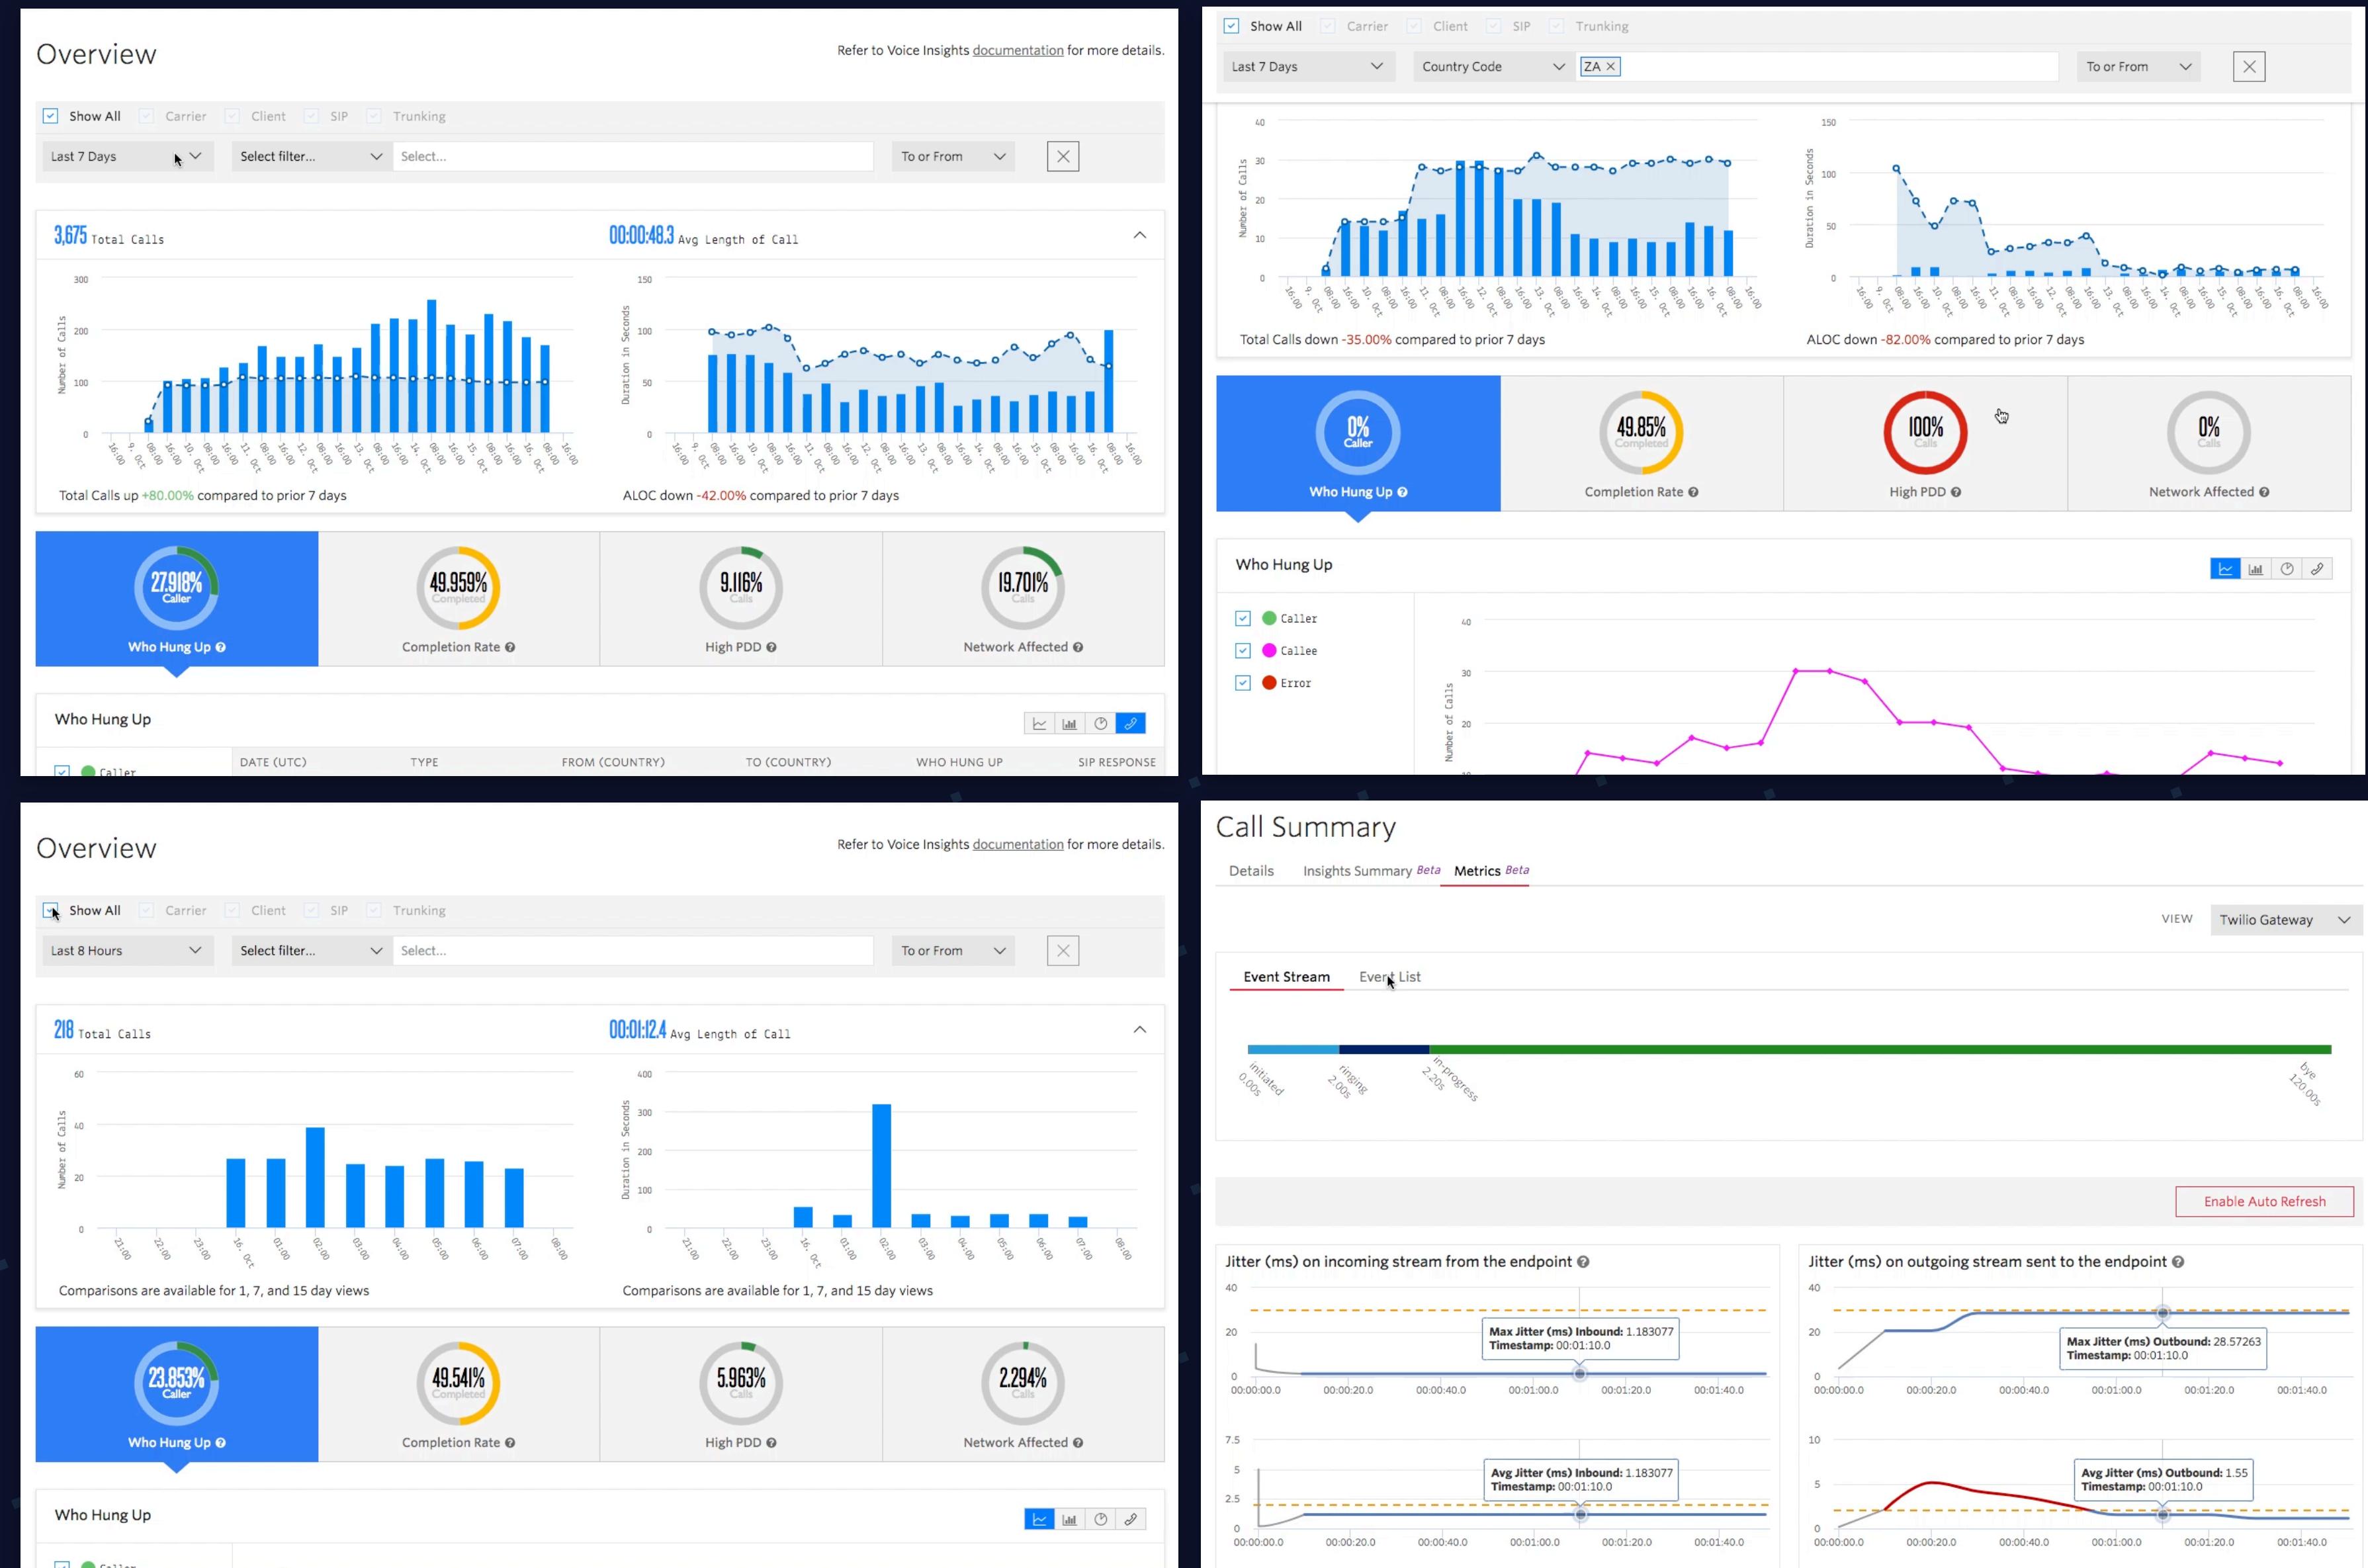
What I Do Voice Insights

Voice Analytics Platform

Millions of complex time series aggregations daily across billions of raw events

Supports internal and external customers.

*Near Real-Time





WORKSHOP

<http://bit.ly/odsc-east-2019-realish-predict>

Data Engineering

Exploratory Data Analysis

Clustering and Tagging

Naive Bayes and Probability

Streaming Predictions



DATA ENGINEERING

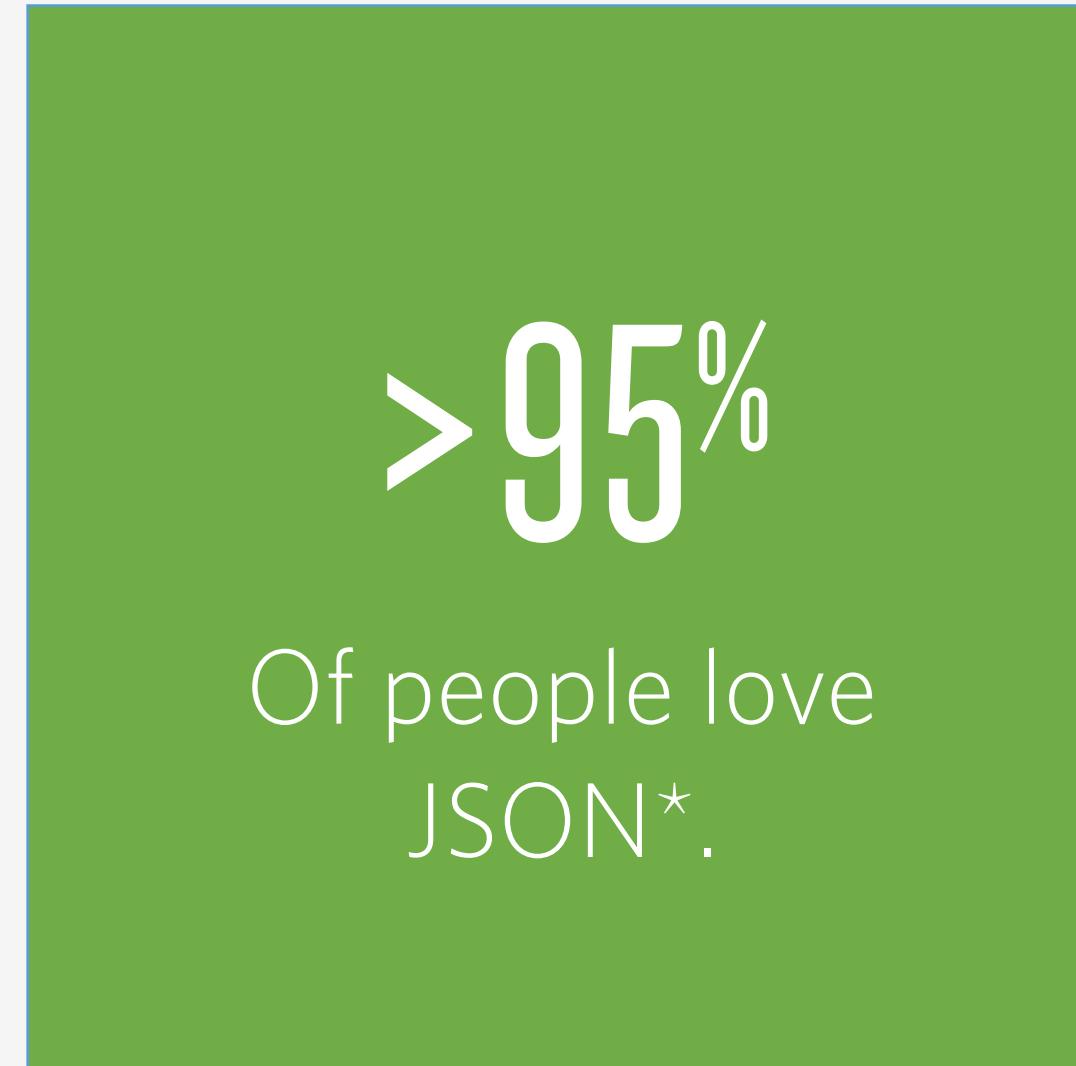
<http://bit.ly/odsc-east-2019-realish-predict>

Data Engineering

- Define Schemas
- ETL (extract transform load)
- Clean and Sanitize Data
- Own your Data

Structured Data

- JSON has structure. But it isn't Structured Data.



>95%

Of people love
JSON*.

```
{  
  "type": "employee",  
  "nickname": "shaines",  
  "attributes": [  
    {  
      "height": 6,  
      "unit": "ft"  
    }  
  ]  
}
```

Structured Data

- JSON has poor runtime guarantees due to its flexible nature. Optimize for compile time guarantees.

```
{  
  "type": "employee",  
  "nickname": null,  
  "attributes": "bae345"  
}
```

100%

Of people saddened
by bad data*

Structured Data

- Protocol Buffers or Avro (arrow, etc) give you **types** and **compile time guarantees**
- **Versioning your API** now just means sticking to a version of your schema.
- Rely on **release for versioning**.
- Perfect for Pipeline interop

```
message Employee {  
    uint64 created = 1;  
    string uuid = 2;  
    string nickname = 3 [default="you"];  
    repeated Attribute attributes = 4;  
}
```

100%

Of people like when
things work between
releases!

Structured Data

- Parquet is awesome format for data at rest
- Columnar Data with Column Stats in partition header
- Plays natively with Apache Spark, Presto, Impala and Hive
- Schema changes across partitions can be tough...

```
spark.read  
  .parquet("hdfs://.../year=2019")
```

100%

Also like not having to define Encoders

ETL

- **Extract** from one source
- **Transform** data or join with other sources
- **Load** into sink (Kafka, HDFS, etc)

```
def readScores(fileName: String): DataFrame = {
    spark.read.schema(rowSchema).json(s"hdfs://localhost:9000/voiceinsights/es/callscores/$fileName")
}

def flattenCallScore(df: DataFrame): DataFrame = {
    val withMetrics = flattenStruct(df, "metrics")
    .drop("metrics").drop("score_data").drop("dimension_hash").drop("model_version").drop("model_type")
    .withColumn("call_sid_enc", base64encode(col("call_sid")))
    .withColumn("account_sid_enc", base64encode(col("account_sid")))
    .drop("call_sid").drop("account_sid")
    .withColumnRenamed("call_sid_enc", "call_sid")
    .withColumnRenamed("account_sid_enc", "account_sid")
}

val withDimensions = flattenStruct(withMetrics, "dimensions").drop("dimensions")
    .withColumn("dimensions_carrier_sid", base64encode(col("dimensions_provider_sid")))
    .drop("dimensions_provider_sid").drop("dimensions_account_sid").drop("dimensions_signaling_ip").drop("dimensions_is_reachable")

expandArray(withDimensions, col("insight_tags")).drop("insight_tags")
    .withColumn("ts", functions.from_utc_timestamp(col("iso_timestamp"), "UTC").drop("iso_timestamp"))
    .withColumn("year", functions.year(col("ts")))
    .withColumn("month", functions.month(col("ts")))
    .withColumn("day", functions.dayofmonth(col("ts")))
    .withColumn("hour", functions.hour(col("ts")))
    .withColumn("dayofyear", functions.dayofyear(col("ts")))
}

def writeScores(df: DataFrame, path: String): Unit = {
    df
    .repartitionByRange(4, col("day"))
    .sortWithinPartitions(col("ts"))
    .write
    .partitionBy("year", "month", "day")
    .mode(SaveMode.Append)
    .format("parquet")
    .option("path", path)
    .save
}
```

Own your Data

- What does your data need to do for you?
- How is storing the data going to help your company or organization?
- Does your data meet GDPR guidelines?





Data Engineering

What we've learned

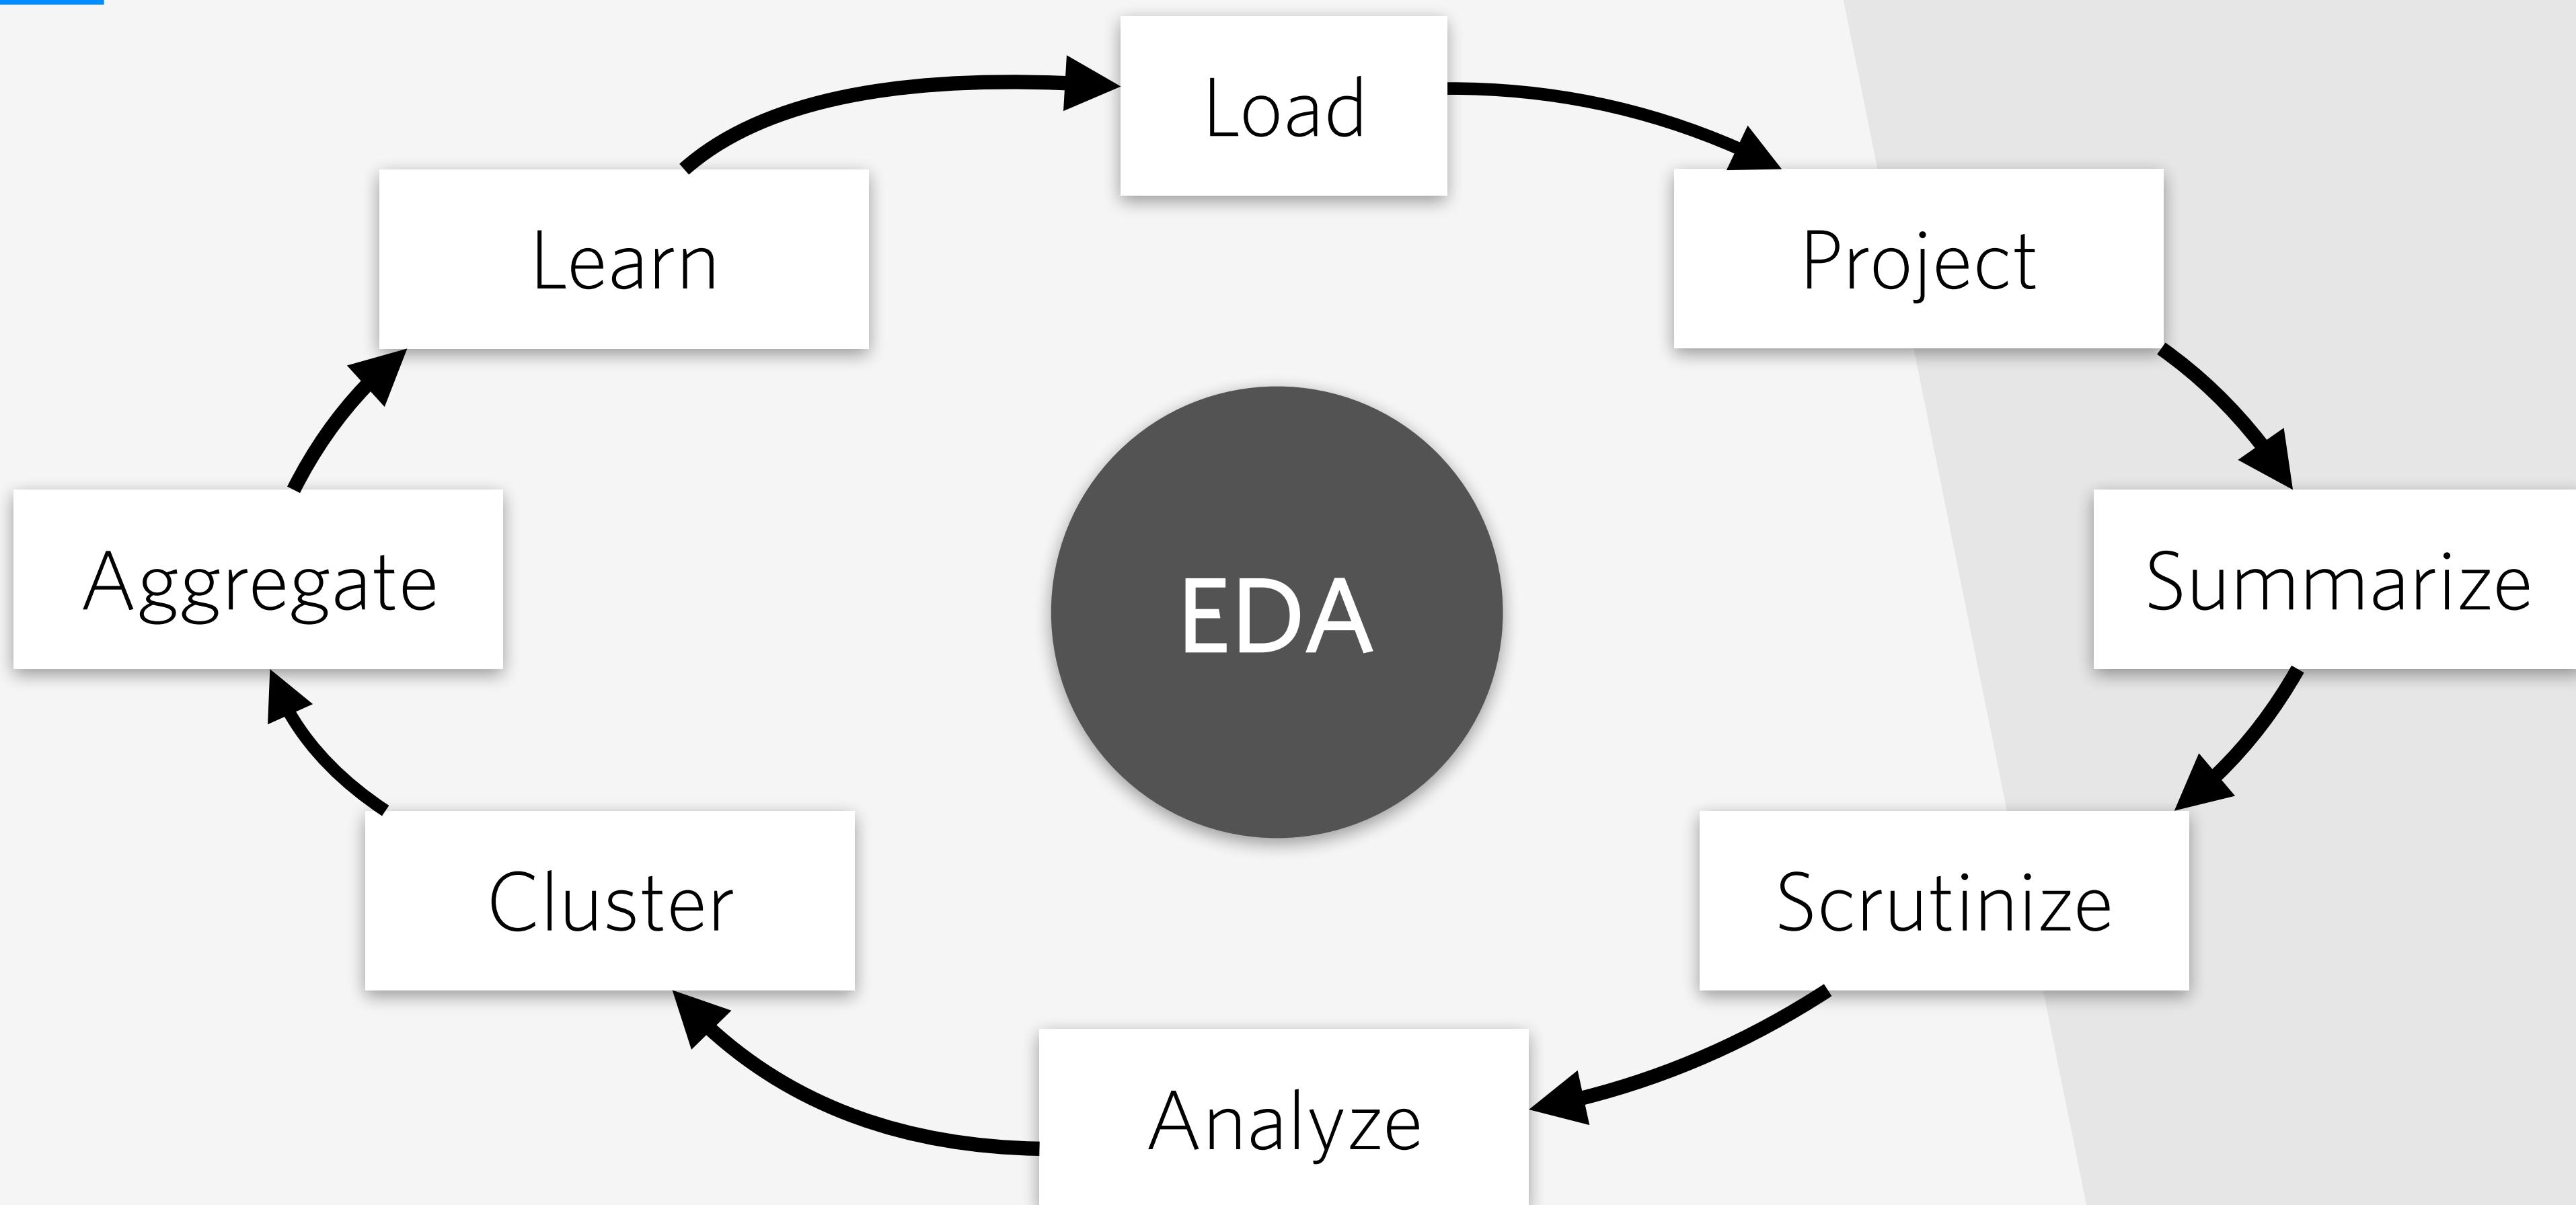
1. **Own** your data. You or your team holds a very valuable asset in your hands and you are the stewards of that Data.
2. **Pipelines** should be built with versioned structured data. **Compile Time guarantees** help keep pipelines healthy.
3. Store only what you will need and **compact your parquet** to reduce FileSystem.io



EXPLORATORY DATA ANALYSIS

<http://bit.ly/odsc-east-2019-realish-predict>

Exploratory Data Analysis



EDA

Zeppelin Notebook ▾ Job Search anonymous ▾

/ODSC-East-2019/GentleIntroduction

import org.apache.spark.sql.functions._
import org.apache.spark.sql.types._
//import com.twilio.open.odsc.realish._

Took 2 sec. Last updated by anonymous at April 30 2019, 7:35:24 PM.

%spark
val scoresData = spark.read.option("basePath", "hdfs://localhost:9000/user/ds/odsc/east
/realish/eda/callscores/").parquet("hdfs://localhost:9000/user/ds/odsc/east/realish/eda

scoresData: org.apache.spark.sql.DataFrame = [timestamp: bigint, score: float ... 33 more fi
elds]

Took 0 sec. Last updated by anonymous at April 30 2019, 7:35:24 PM. (outdated)

%spark
scoresData.inputFiles

Took 1 sec. Last updated by anonymous at April 30 2019, 7:35:25 PM. (outdated)

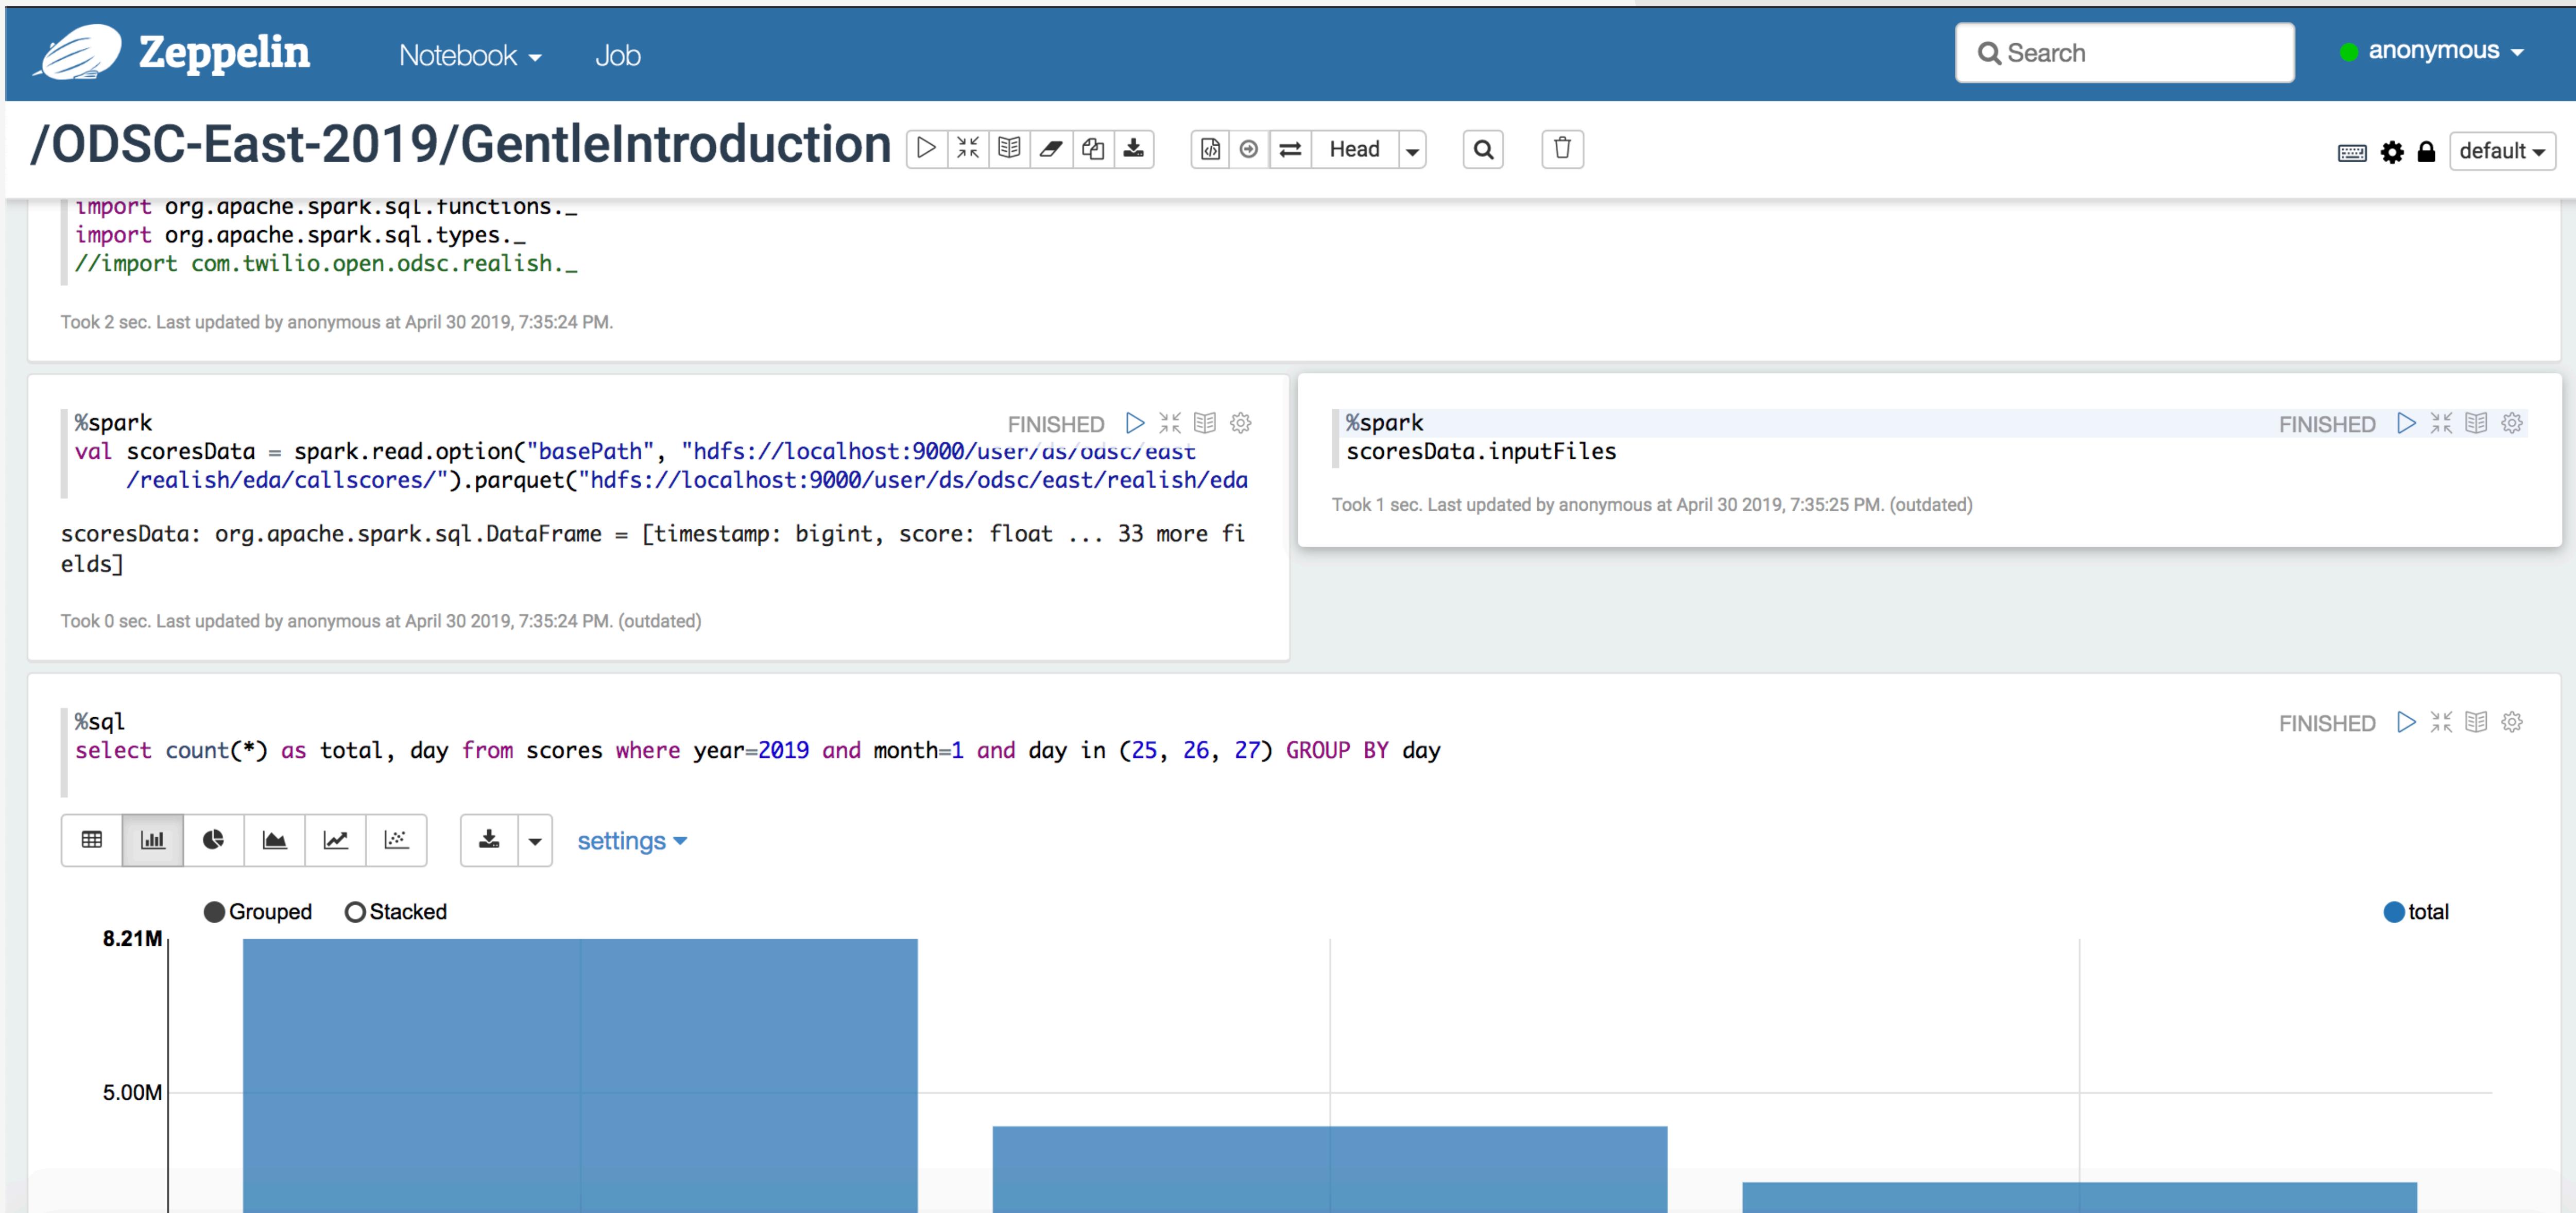
%sql
select count(*) as total, day from scores where year=2019 and month=1 and day in (25, 26, 27) GROUP BY day

FINISHED

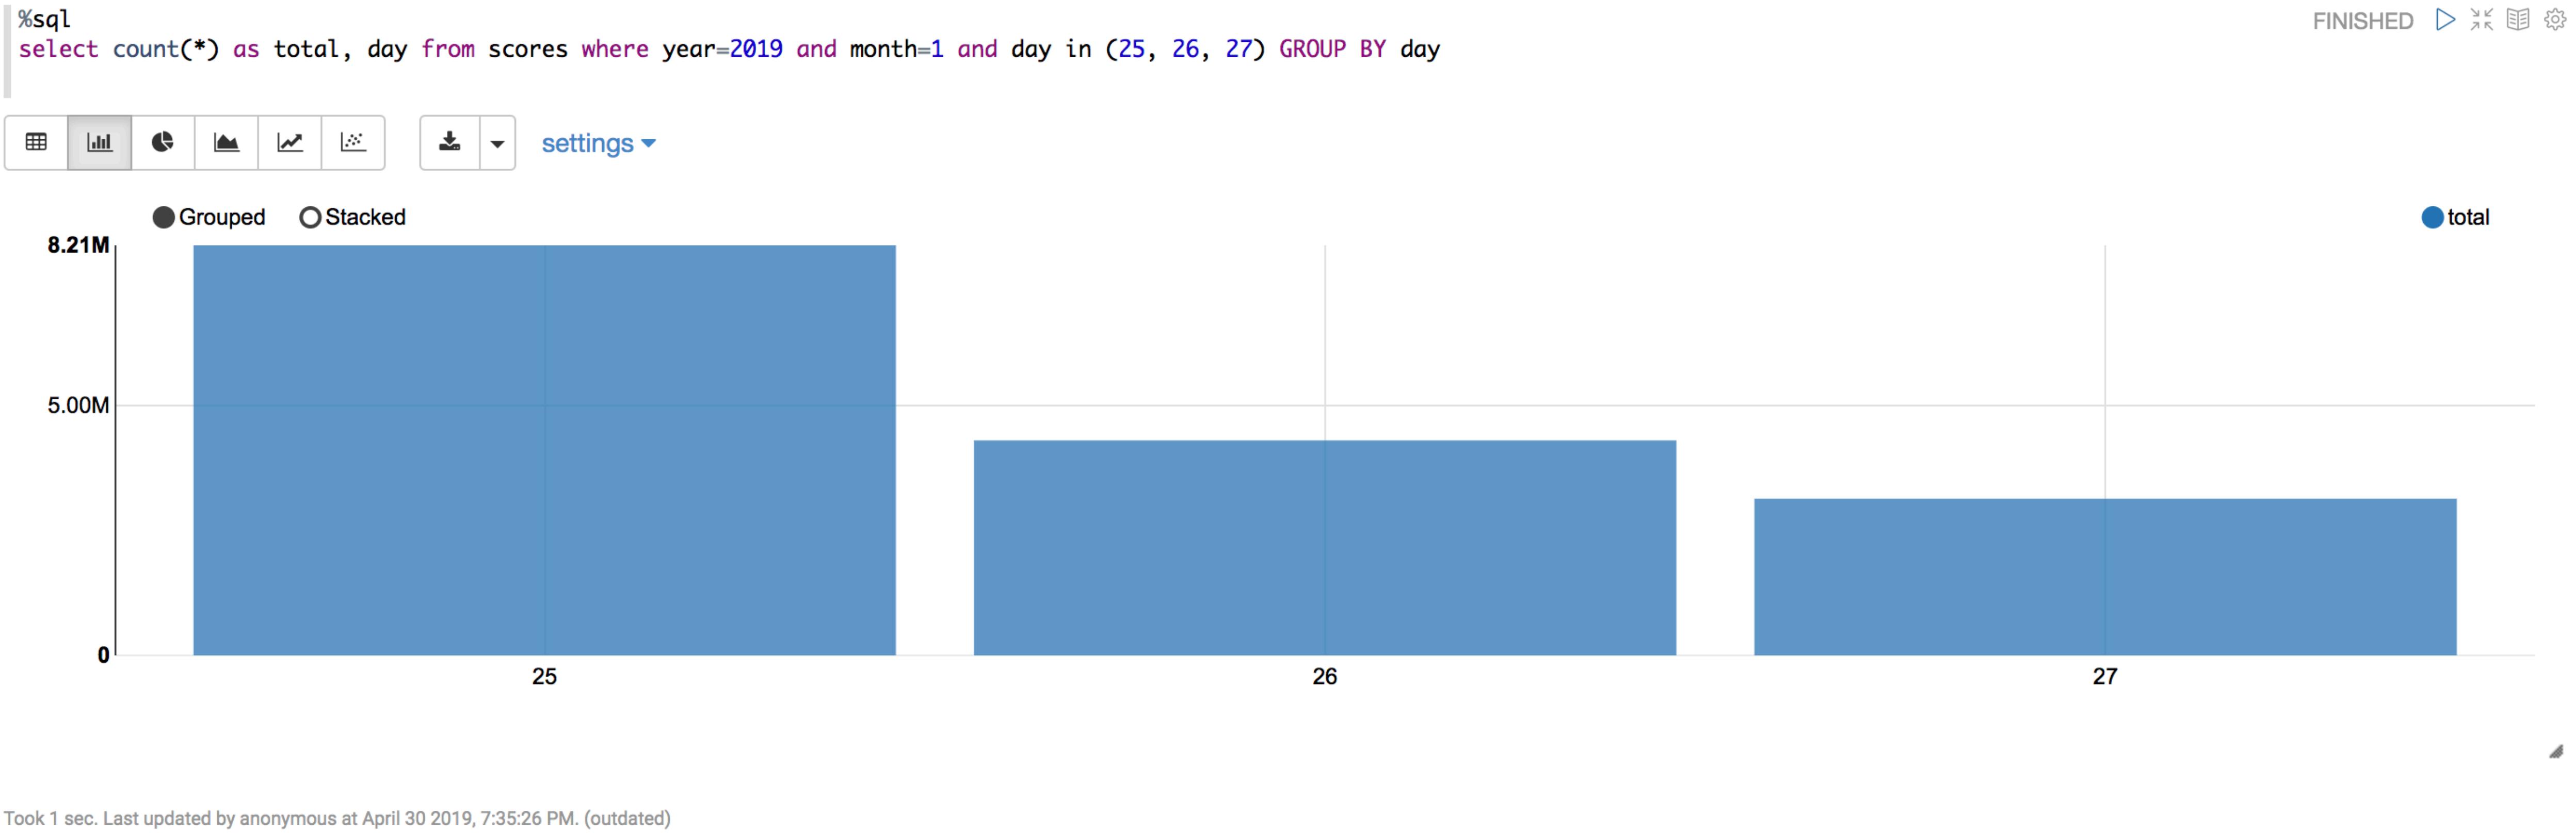
Grouped Stacked settings

8.21M
5.00M

total



EDA



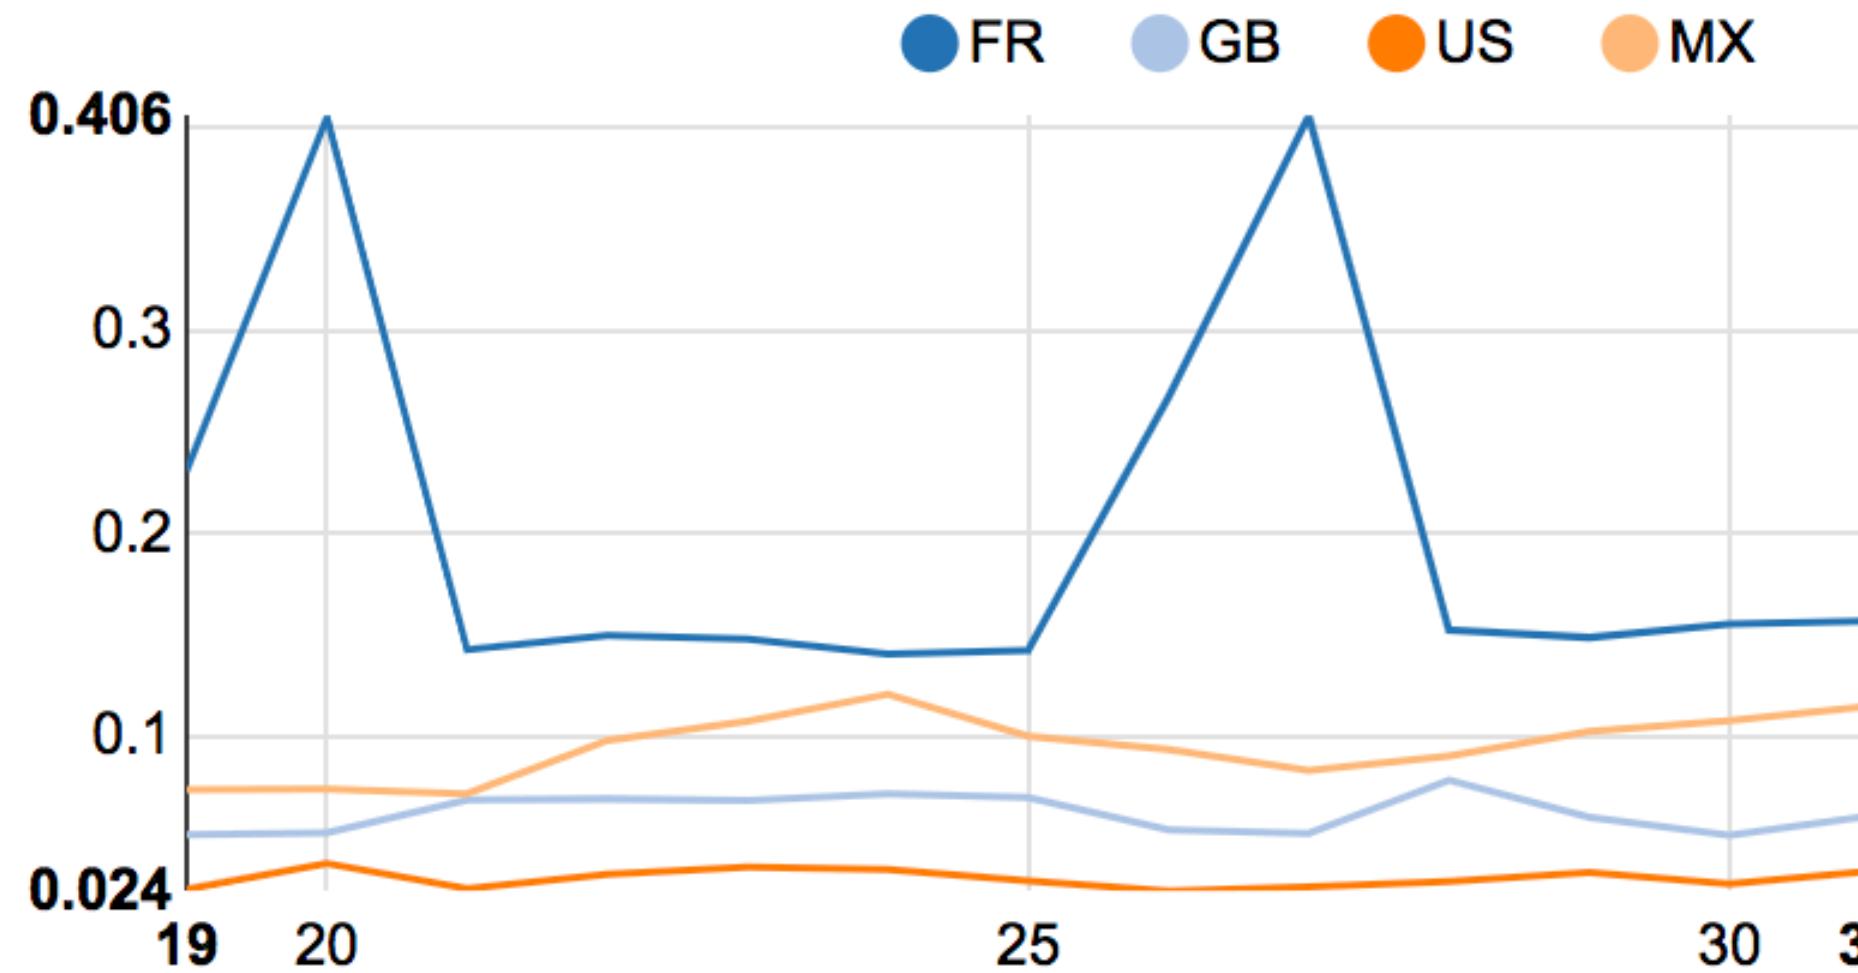
https://github.com/newfront/odsc-east-realish-predictions/blob/master/zeppelin/odsc_east_intro_part1.json

EDA

```
%sql
select dayofyear, dimensions_country_code, avg(pktloss_sum
/total) as avg_pktloss_percentage from
carrier_by_country_scores where
dimensions_country_code in ("US", "FR", "GB", "MX")
```



settings ▾

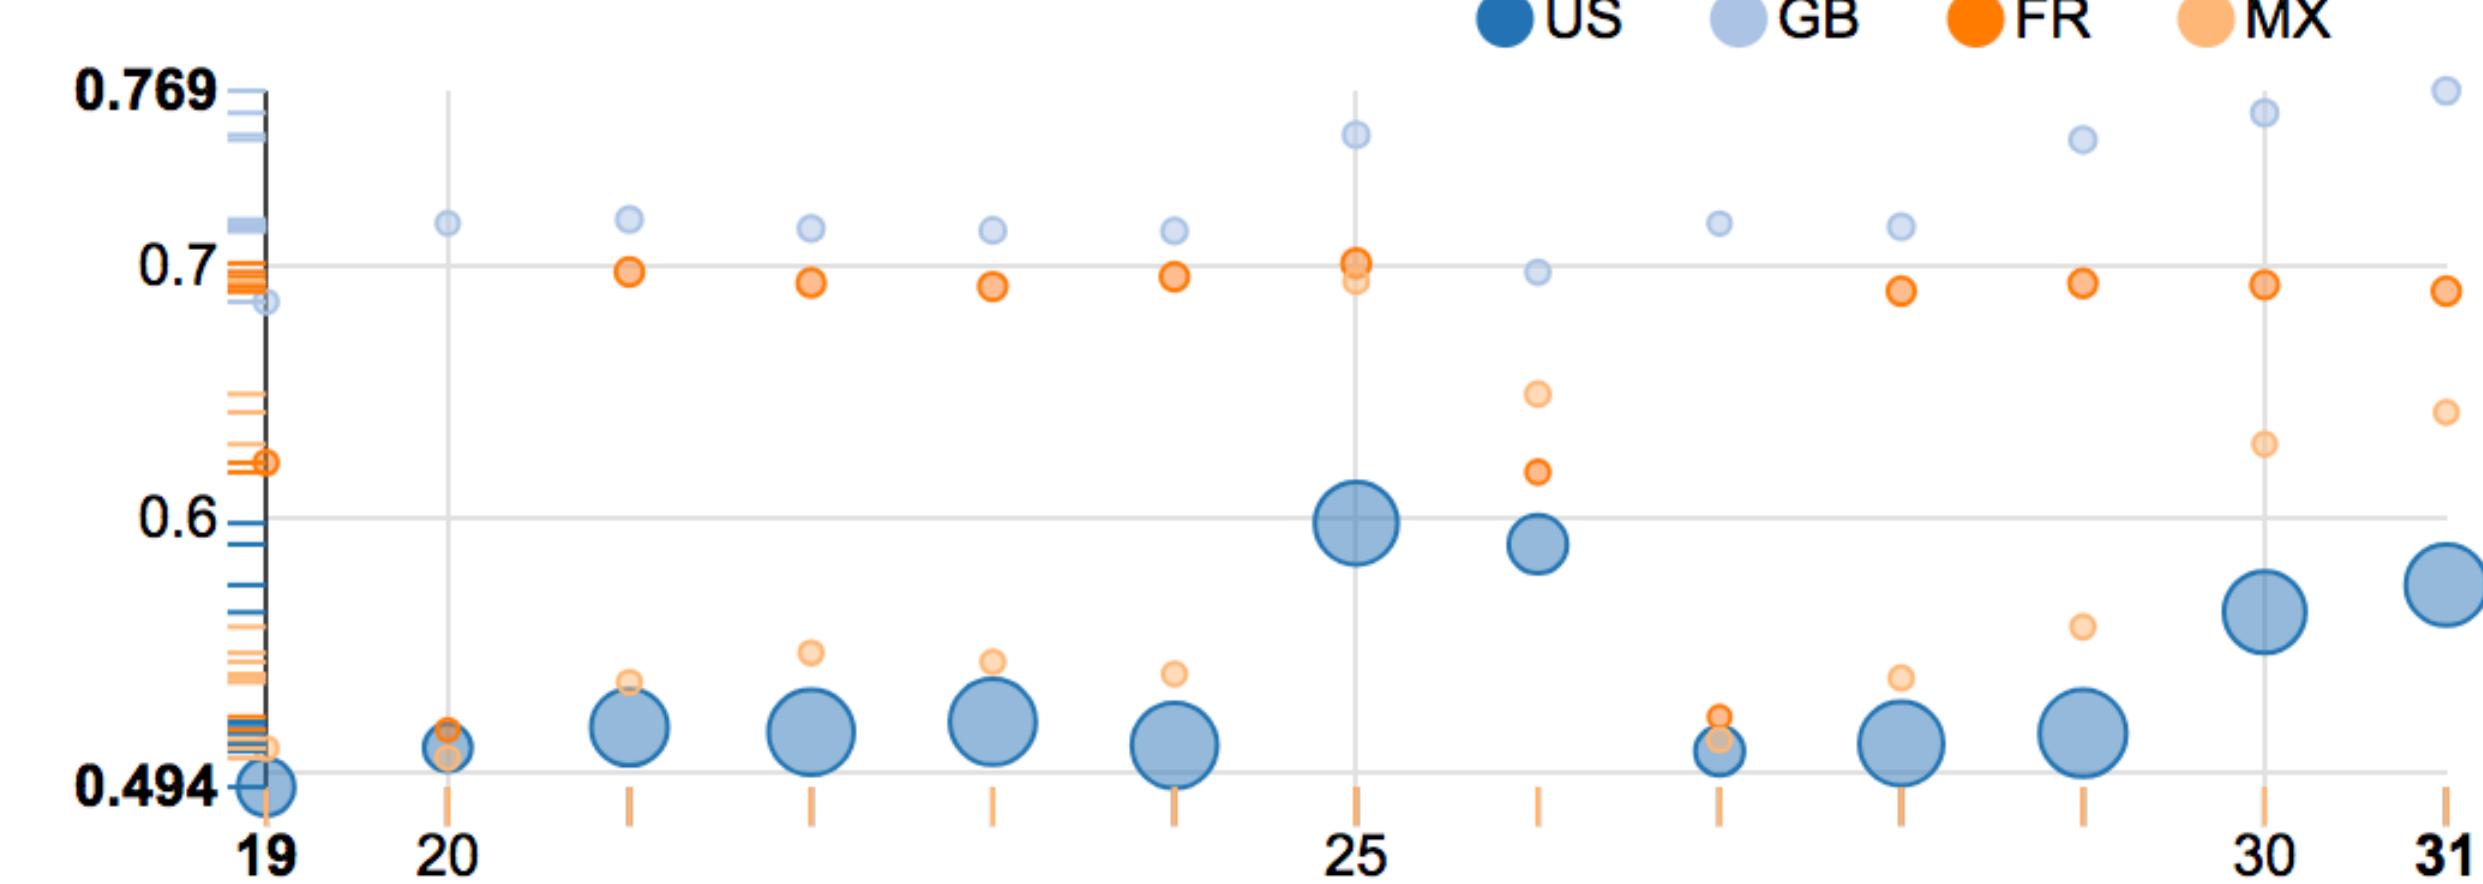


Took 9 sec. Last updated by anonymous at April 30 2019, 7:35:59 PM. (outdated)

```
%sql
select count(*) as total, min(score) as min_score, avg(score) as
mean_score, max(score) as max_score, dimensions_country_code,
dayofyear from scores where year=2019 and month=1 and
dimensions_country_code IN ("US", "FR", "GB", "MX") GROUP BY
dimensions_country_code, dayofyear HAVING total > 1000 ORDER BY
```



settings ▾



Took 8 sec. Last updated by anonymous at April 30 2019, 7:36:07 PM.

EDA

Zeppelin Notebook ▾ Job Search anonymous ▾

ODSC-East-2019/ExploratoryDataAnalysis

%spark
val scoresData = spark.read.option("basePath", "hdfs://localhost:9000/user/ds/odsc/east/realish/eda/callscores/").parquet("hdfs://localhost:9000/user/ds/odsc/east/realish/eda/callscores/")
scoresData.createOrReplaceTempView("scores")

scoresData: org.apache.spark.sql.DataFrame = [timestamp: bigint, score: float ... 33 more fields]

SPARK JOB FINISHED Took 1 sec. Last updated by anonymous at May 01 2019, 1:13:16 AM.

%sql
describe scores

col_name	data_type	comment
timestamp	bigint	null
score	float	null
metrics_connect_duration	int	null
metrics_pdd	float	null
metrics_avg_jitter	float	null
metrics_max_jitter	float	null
metrics_avg_latency	float	null
metrics_max_latency	float	null

FINISHED Took 0 sec. Last updated by anonymous at April 30 2019, 6:34:56 PM. (outdated)

%spark
DataFrameUtils.distinctColumnValues(scoresData.select("day").where(col("month").equalTo(1))).(spark).show

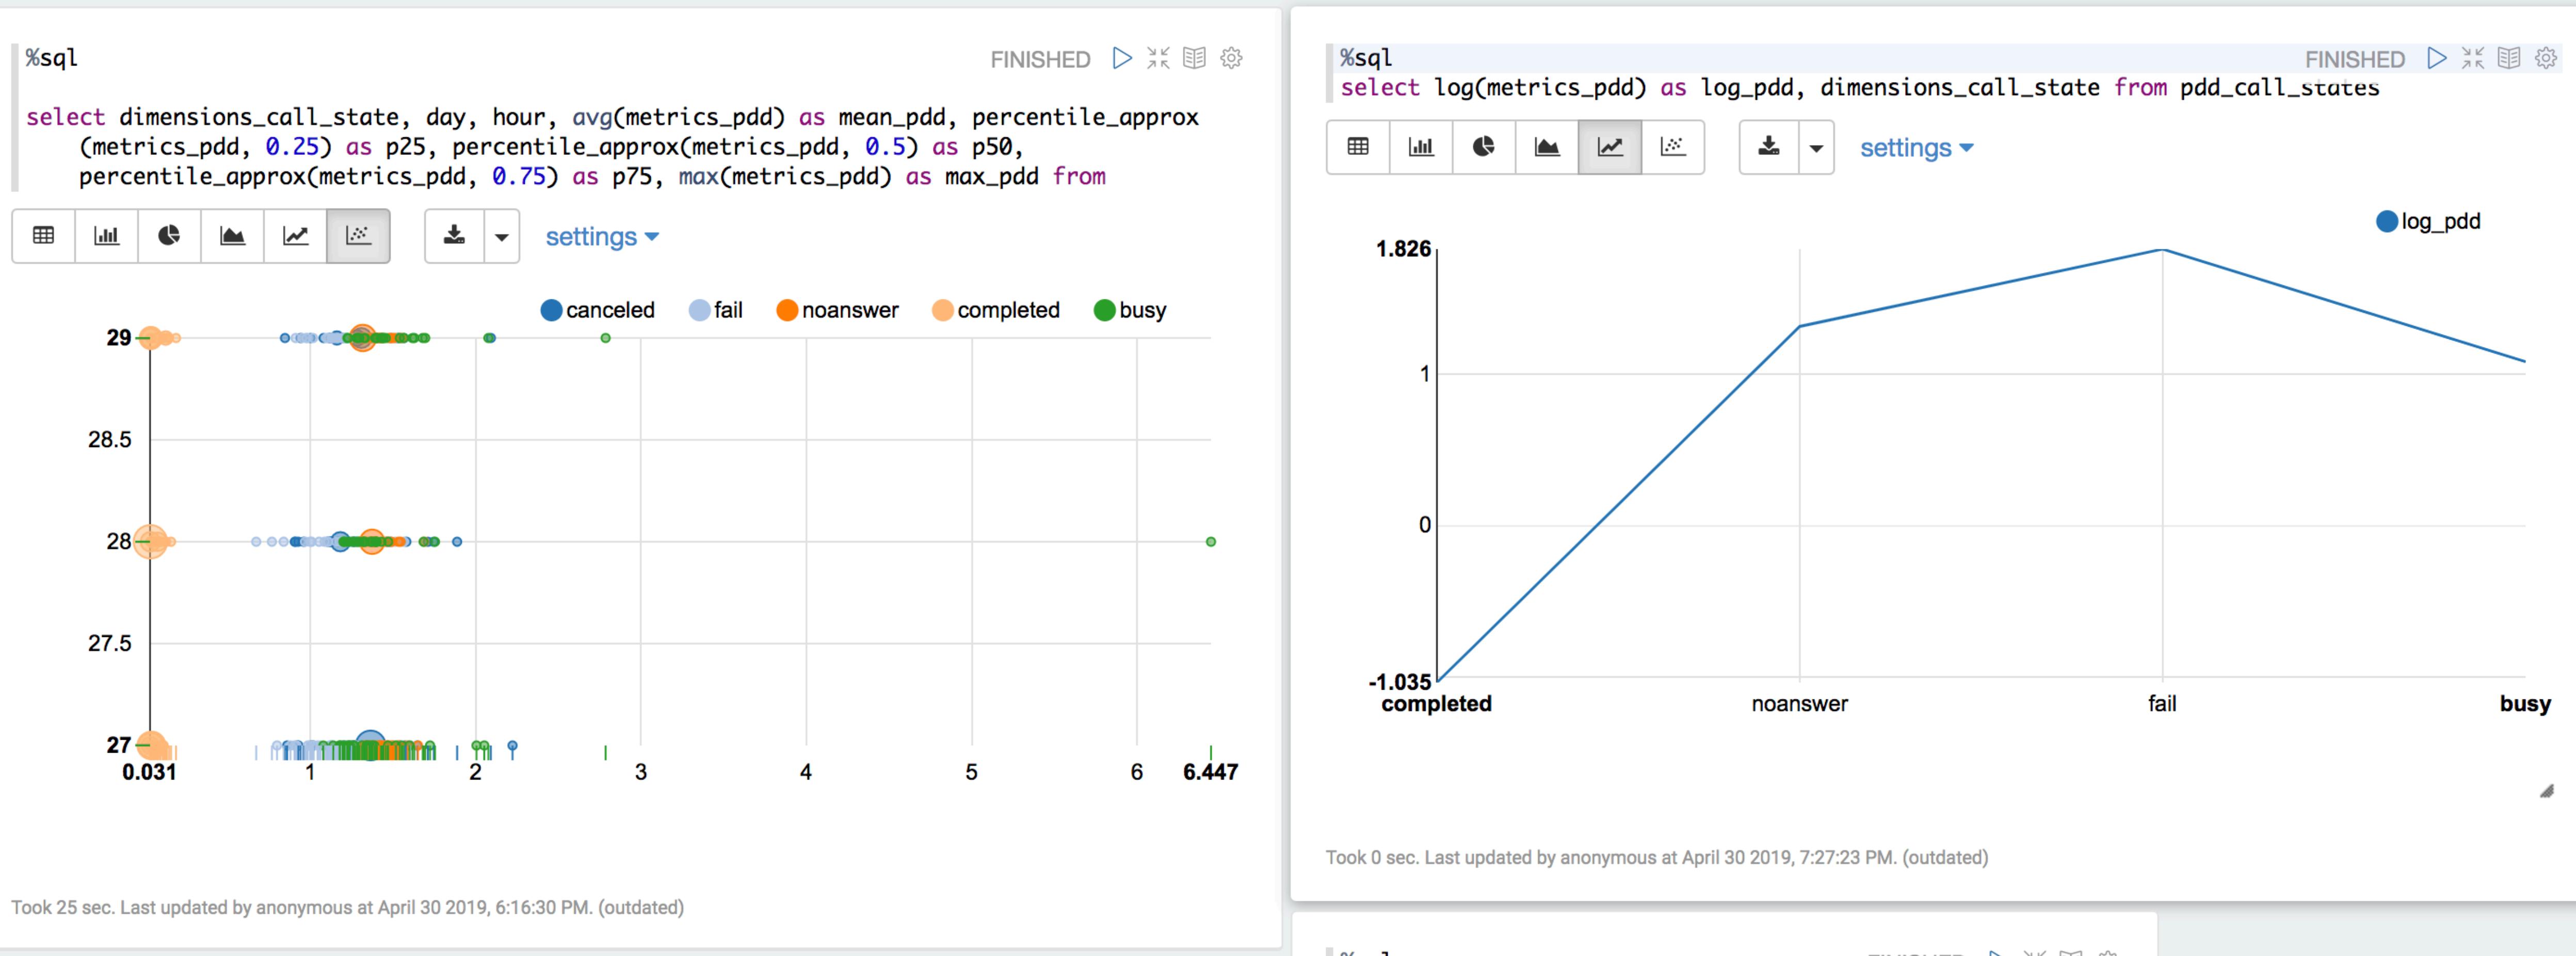
day
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31

SPARK JOB FINISHED Took 2 sec. Last updated by anonymous at May 01 2019, 1:13:26 AM.

%sql
select count(*) as total, day from scores where year=2019 and month=1 group by day order by day

total	day
8035457	31
7989784	30
8846214	29
8420984	28

EDA



https://github.com/newfront/odsc-east-realish-predictions/blob/master/zeppelin/odsc_east_eda_part2.json



Exploratory Data Analysis

What we've learned

1. Uncover the **Shape of your Data**
2. **Find interesting relationships** between values in your data
3. Understand what is Normal.



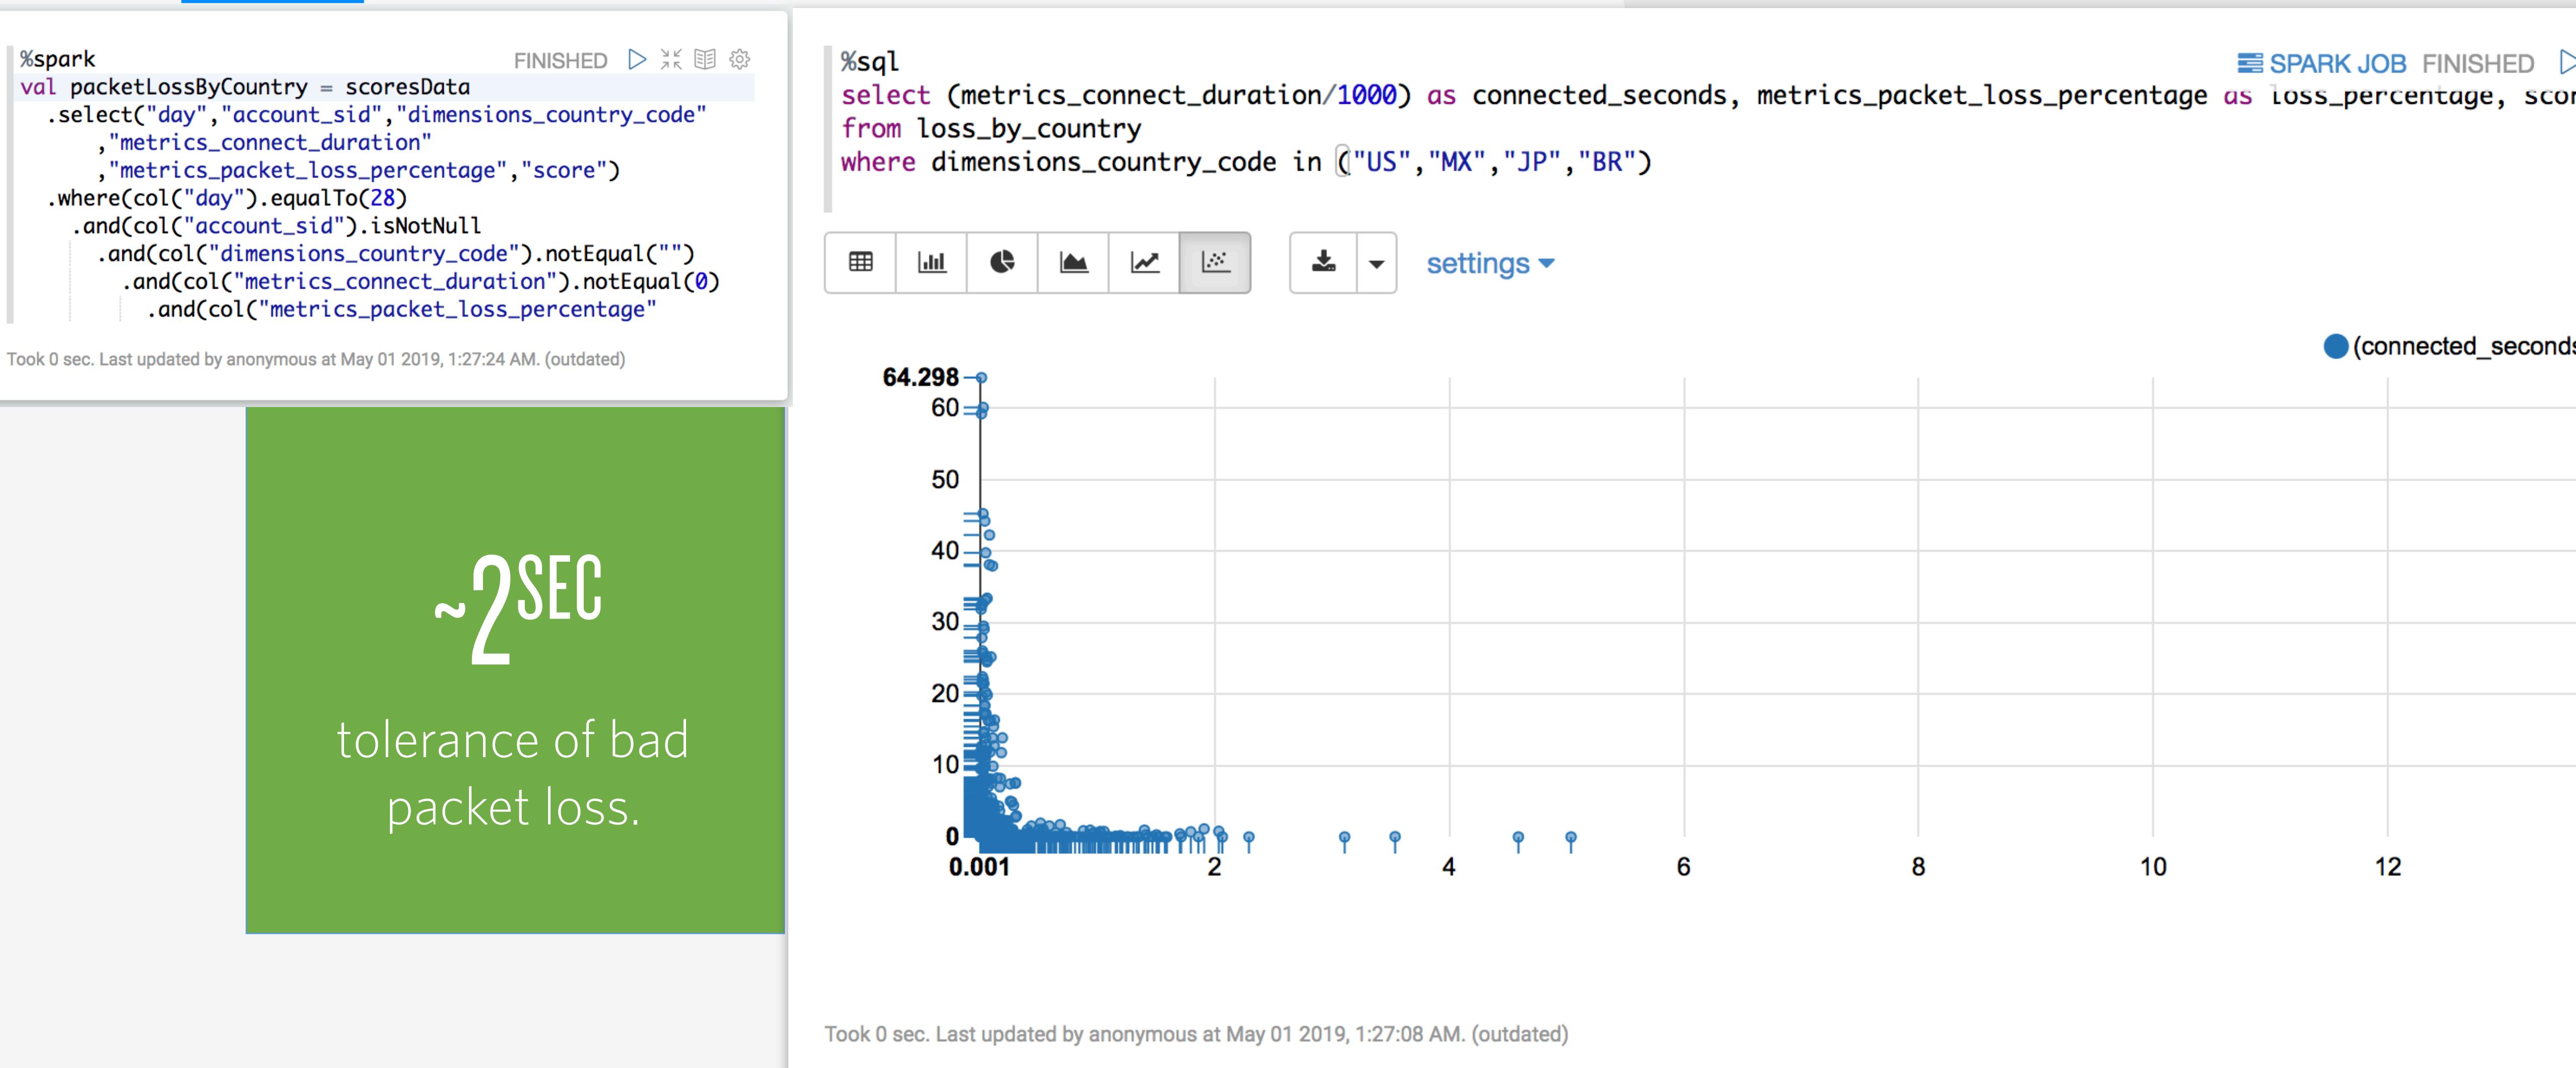
CLUSTERING AND TAGGING

<http://bit.ly/odsc-east-2019-realish-predict>

Clustering and Tagging

- Look at the Universal Set vs SubSet Trends
- Find Patterns in the Noise
- Understand where natural thresholds reside

Clustering and Tagging





Clustering and Tagging

What we've learned

1. Simple **GroupBy** and **Filtering** yields Insight into trends in your data.
2. **Applying tags** at specific thresholds in your data will allow you to find signal in the noise.



STREAMING PREDICTIONS

<http://bit.ly/odsc-east-2019-realish-predict>

Streaming Predictions

```
val streamingQuery = metricData.toDF()
    .withWatermark("timestamp", "2 hours")
    .groupBy(col("metric"), col("countryCode"), window($"timestamp", "5 minutes"))
    .agg(
        min("value") as "min",
        avg("value") as "mean",
        max("value") as "max",
        count(*) as "total"
    )
    .writeStream
    .format("memory")
    .queryName("datastream")
    .outputMode(OutputMode.Append())
    .trigger(processingTimeTrigger)
    .start()

metricData.addData(backingData)

streamingQuery.processAllAvailable()

spark.sql("select * from datastream").show(20, false)

val checkChange = spark.sql("select * from datastream")
    .groupBy("metric", "countryCode")
    .agg(
        sum("total") as "total",
        avg("mean") as "mean"
    )

checkChange.show(20, false)

// now can do interesting things with minor back tracking...
```



Streaming Predictions

What we've learned

1. Given a simple streaming application. You can now go off and build a masterpiece!
2. Remember. Keep your record sizes small or you'll have to cough up a ton of CPU and RAM to handle those poor choices. We did!

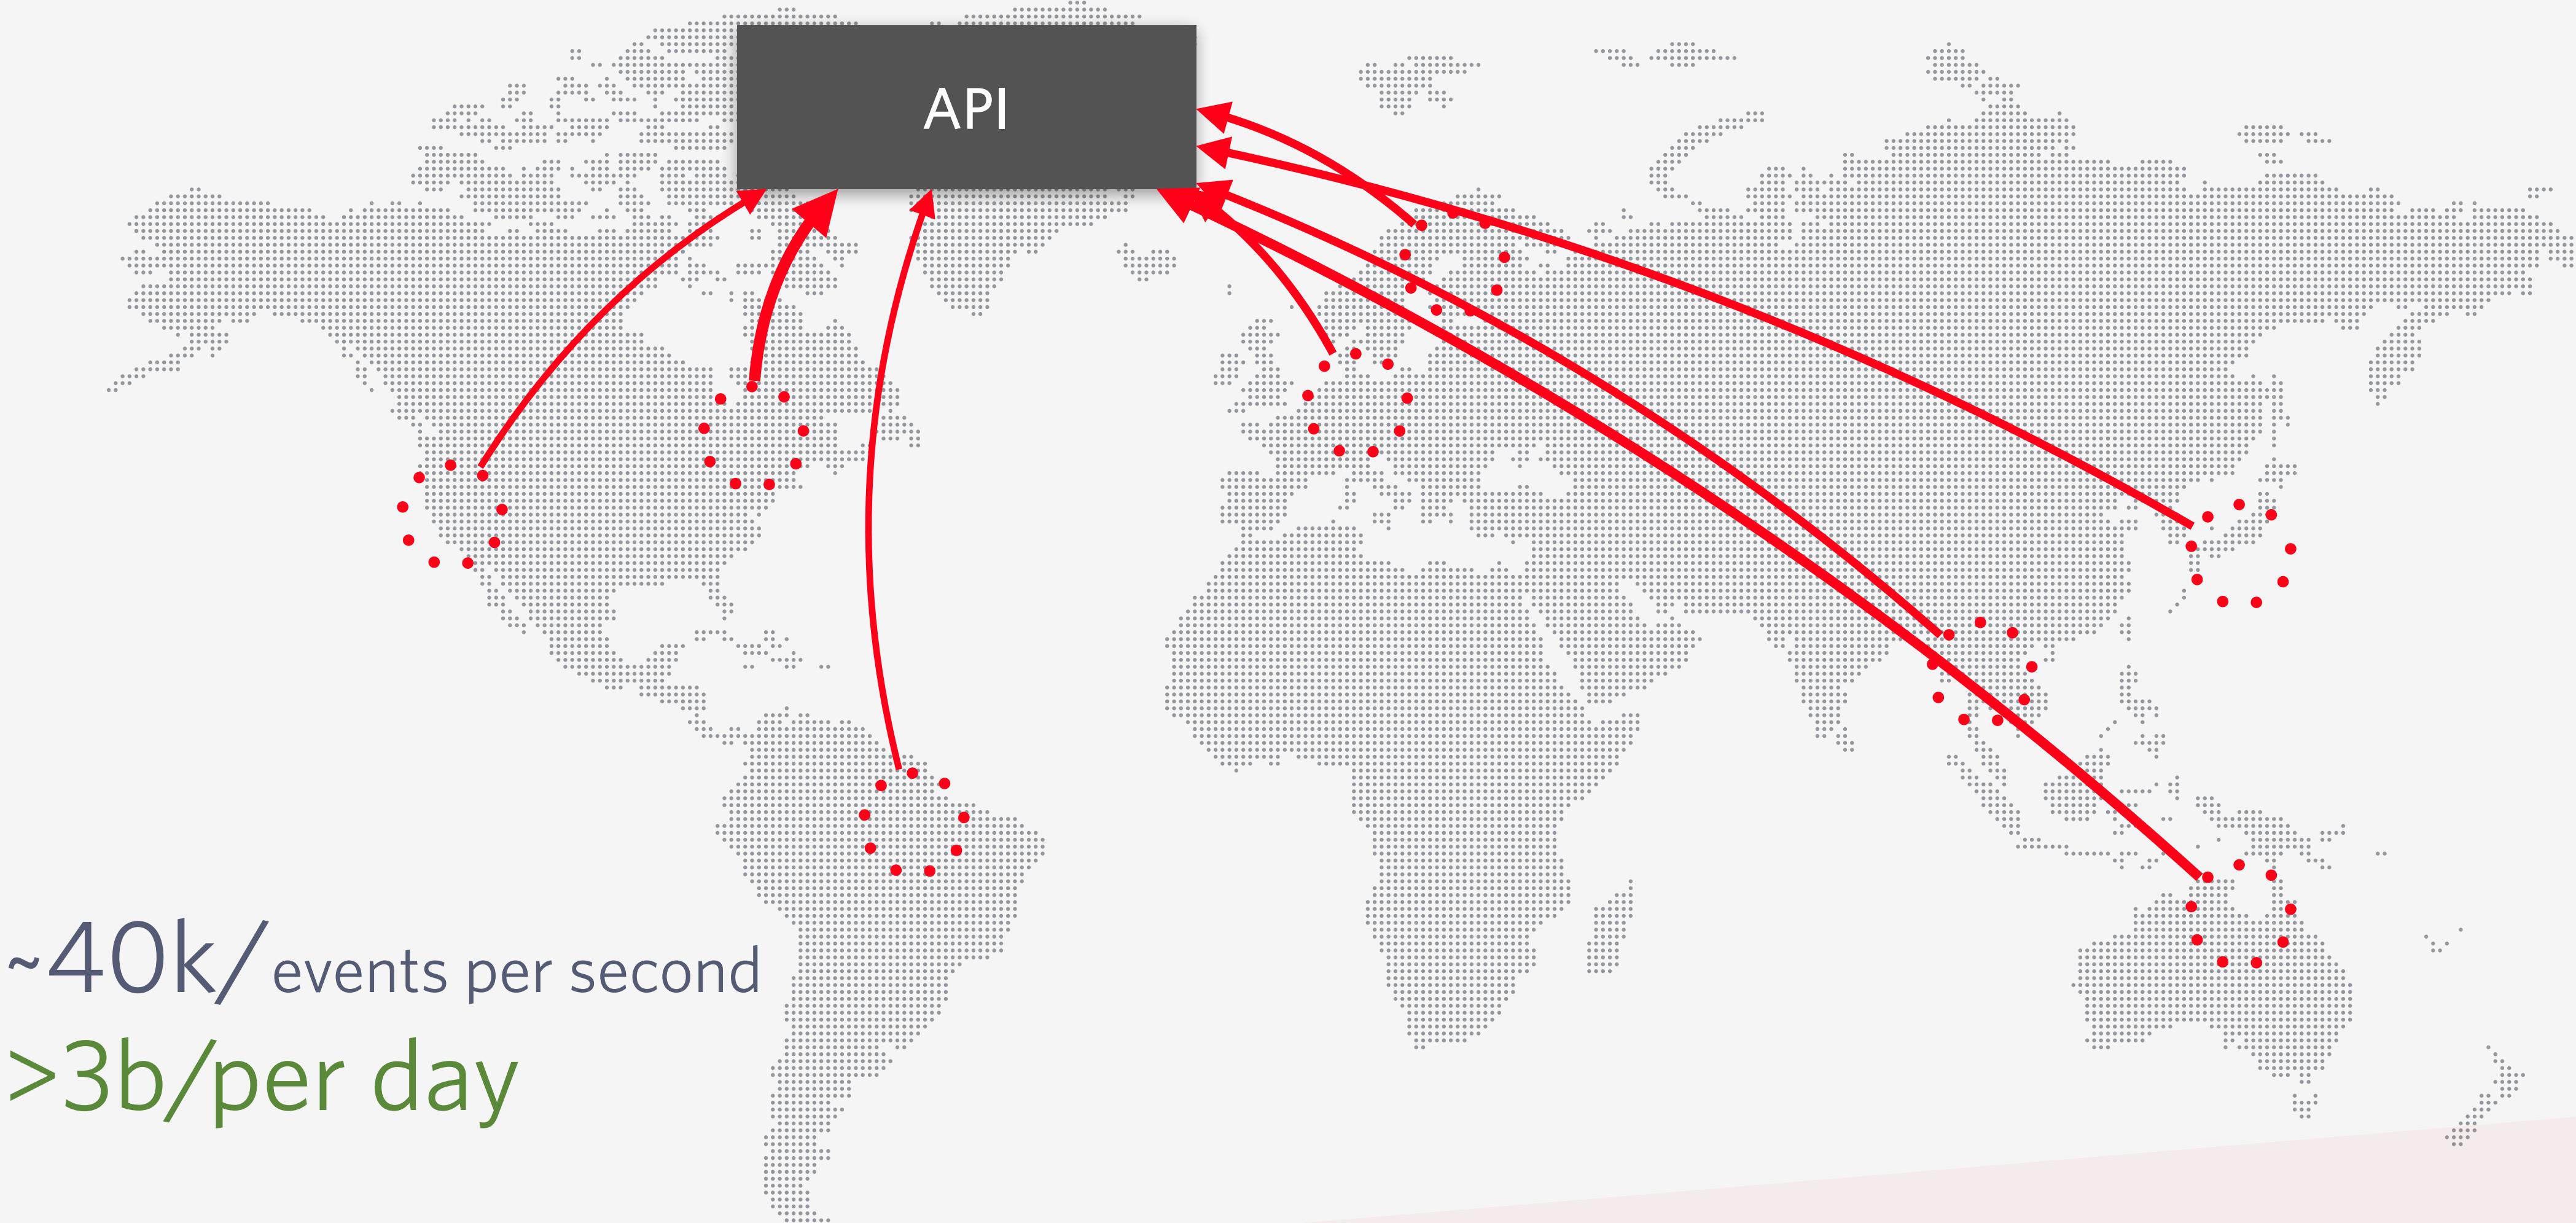
DISTRIBUTED FIRST ARCHITECTURE





DISTRIBUTED FIRST ARCHITECTURE

Why?





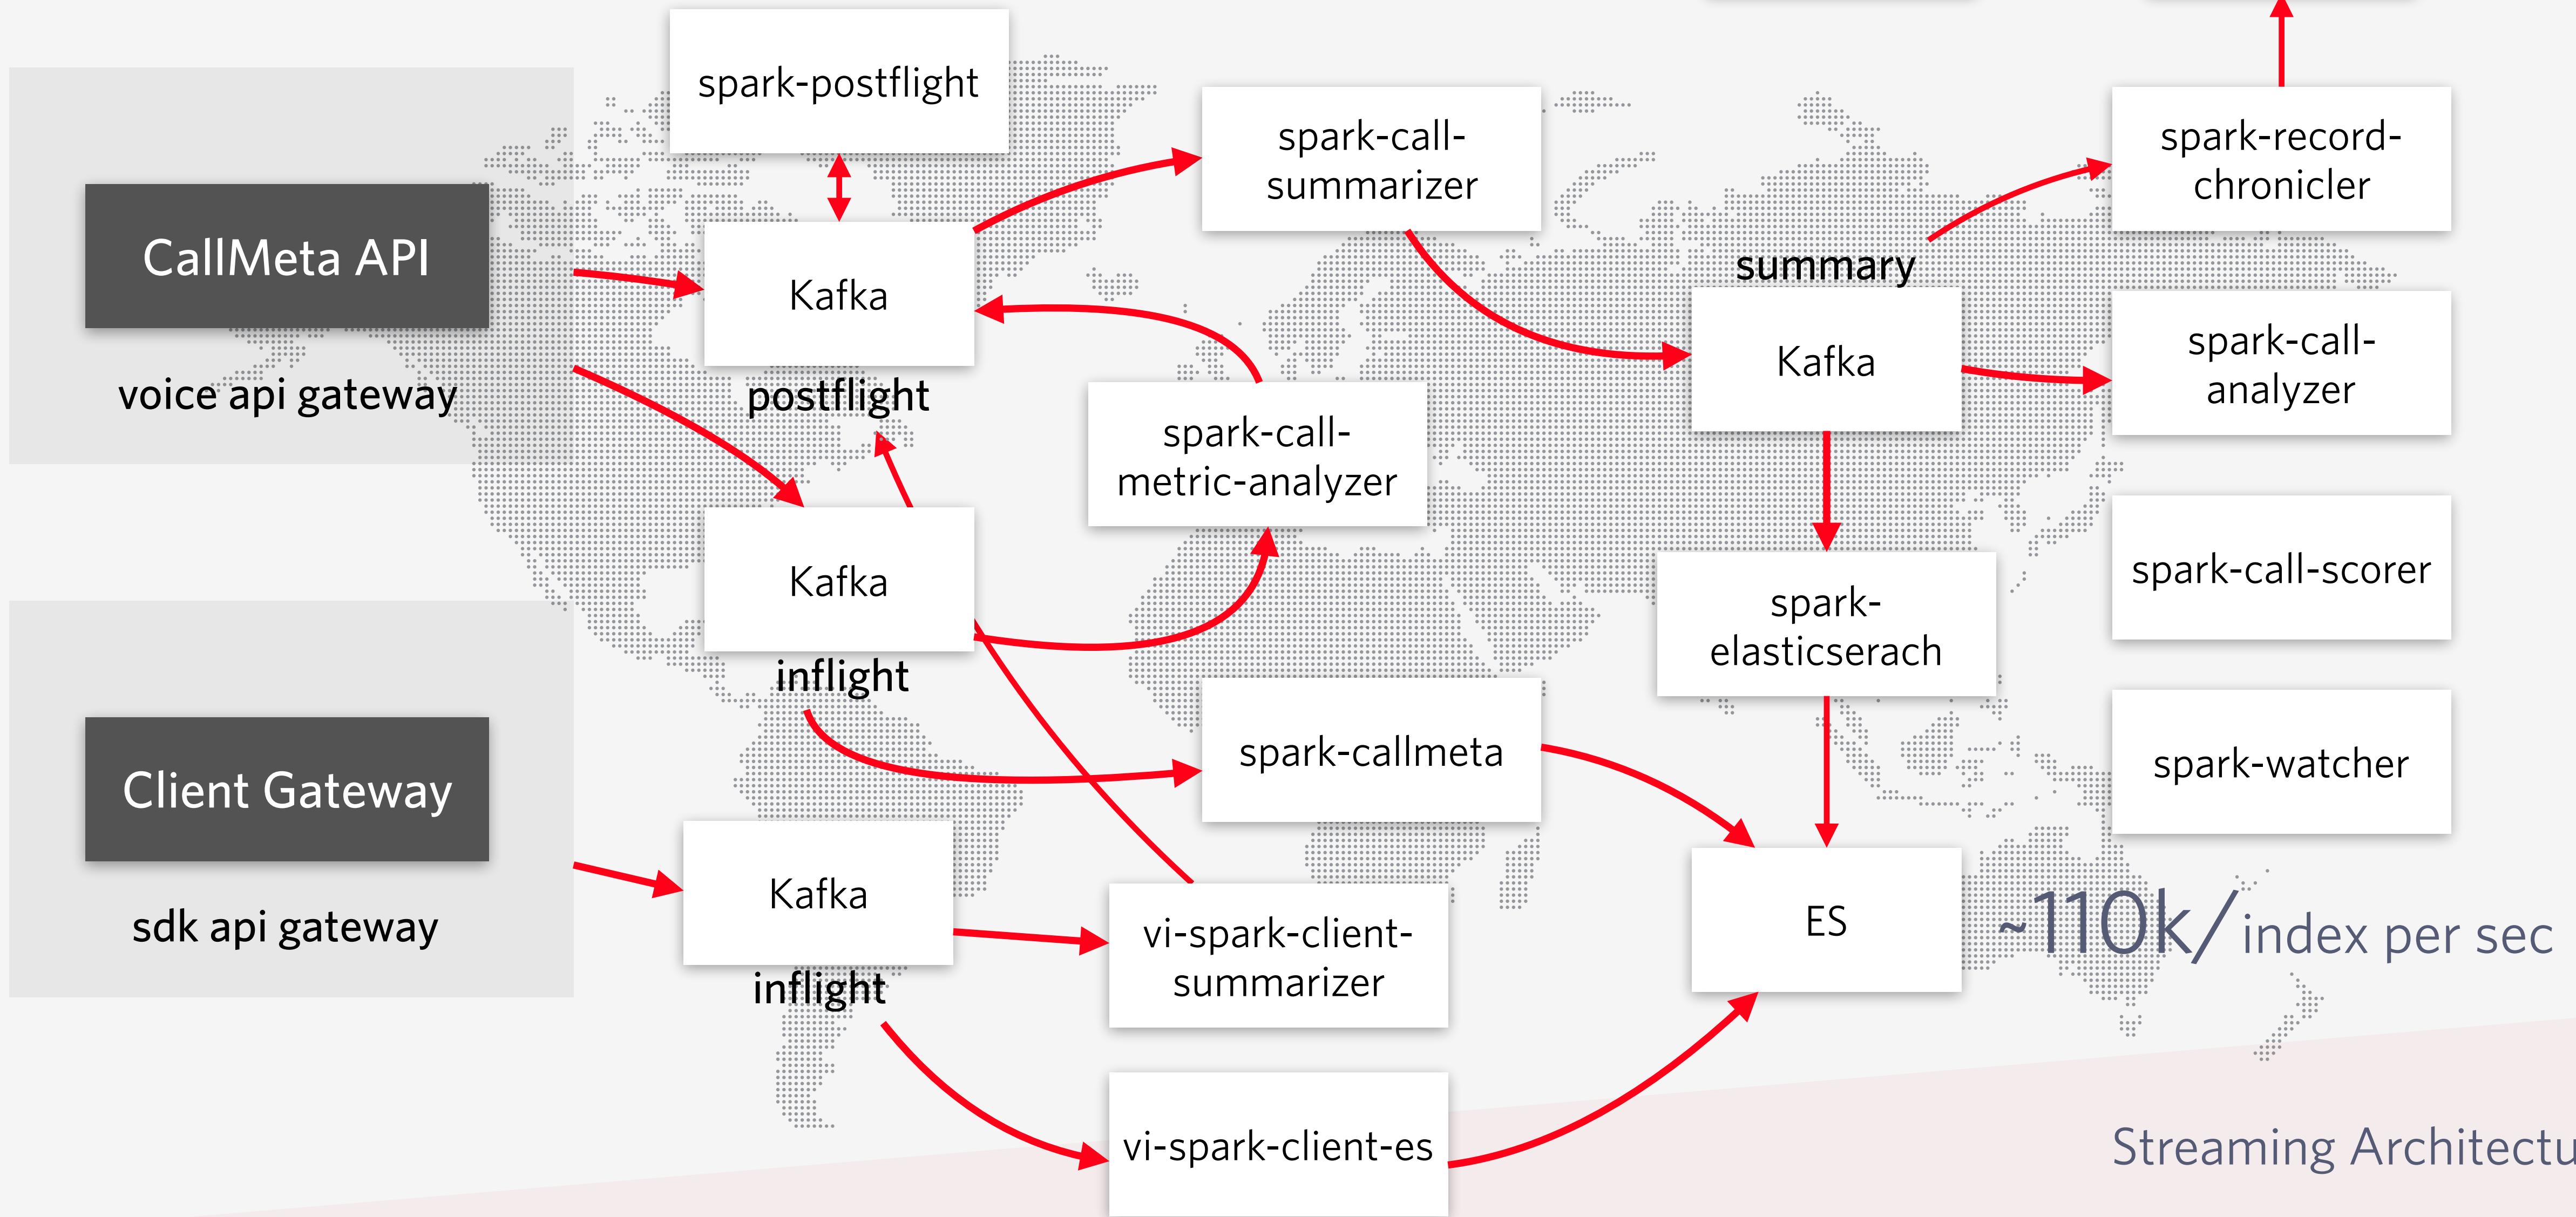
DISTRIBUTED FIRST ARCHITECTURE

How





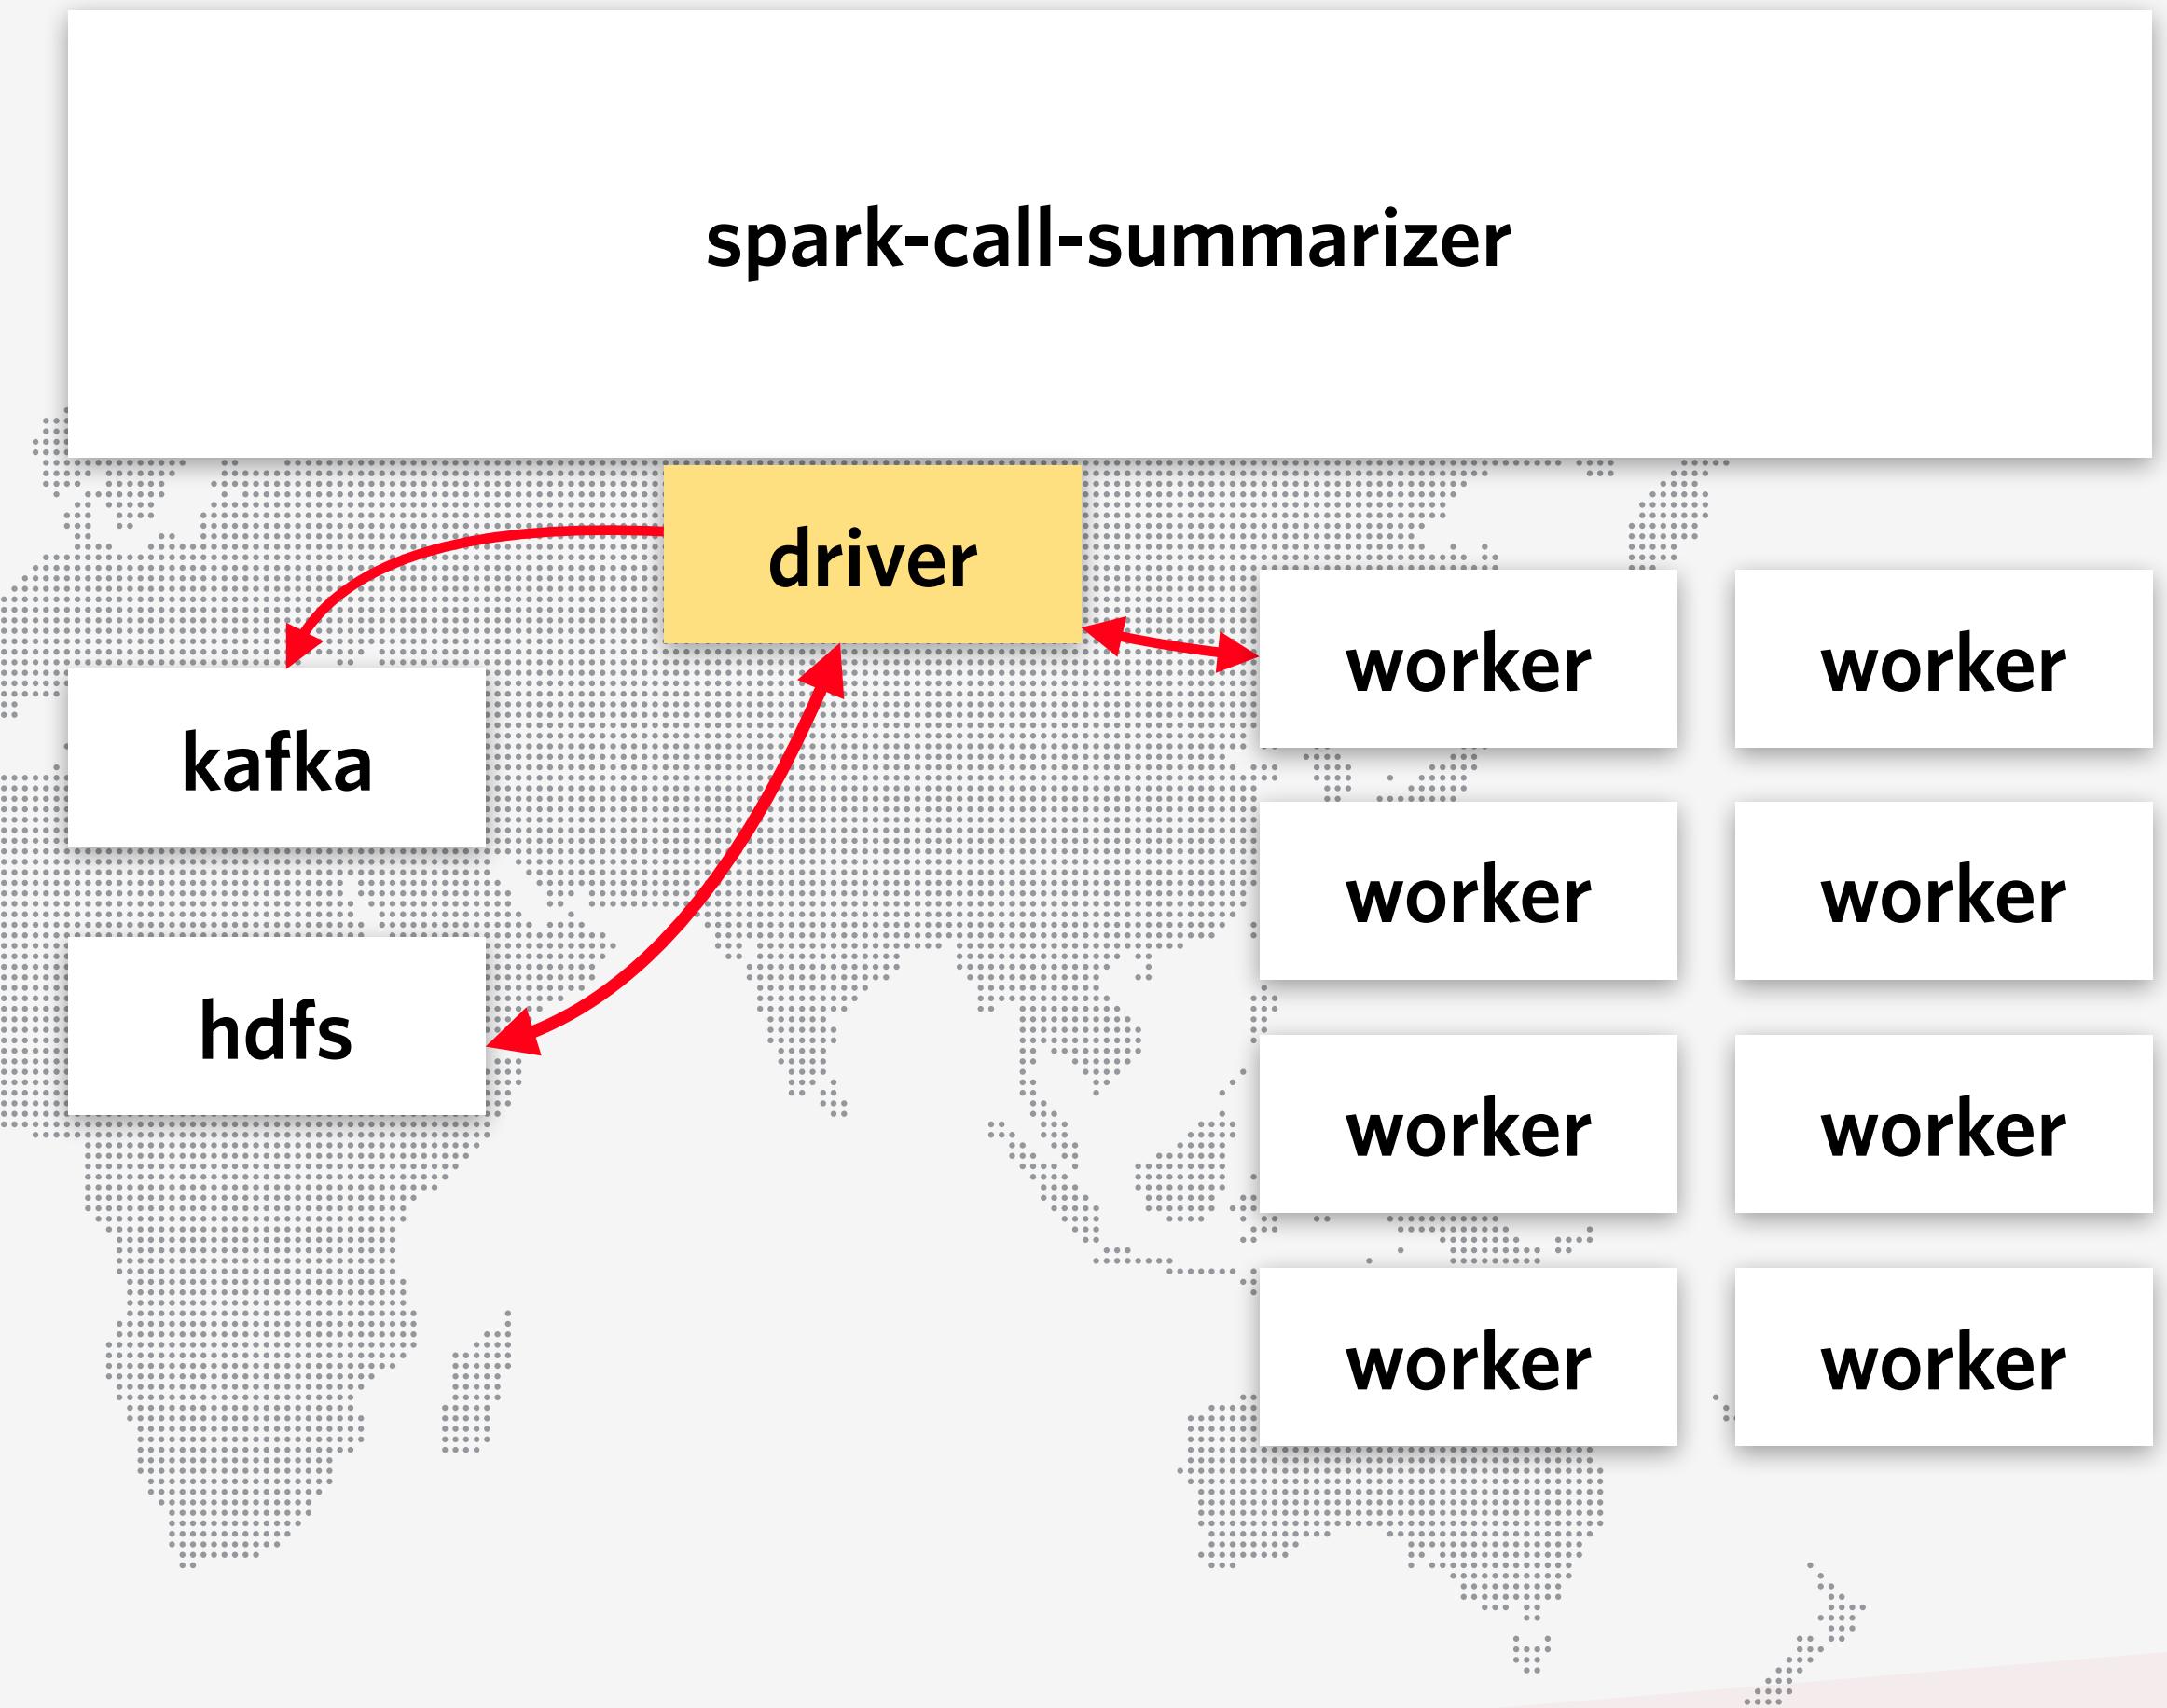
DISTRIBUTED FIRST ARCHITECTURE





SPARK APP ARCHITECTURE

- Each App **does one thing well.**
- Apps can be **deployed via Admiral.**
- Apps can be scaled up or down independently.
- Apps can be run in recovery or real-time mode
- Apps **checkpoint state to HDFS and state survive between releases and upgrades**



Streaming Architecture



SPARK APP ARCHITECTURE

- Apps can be deployed to one or many logical spark clusters

Reload this page

voice-spark-driver

MANIFEST FORMATION

MANIFEST	FORMATION
spark-call-analyzer-034-016	cluster_240_eden
spark-call-scorer-020-015	cluster_240_eden_scorer
spark-super-summarizer	cluster_231_eden
voice-spark-record-chronicler-016	cluster_240_eden
spark-call-watcher-007-002	cluster_231_eden_watcher

ASSIGN MANIFEST ASSIGN CONFIGURATION COPY DELETE

Streaming Architecture

twitter: [@newfront](#)
blog: blog.twilio.com

THANK YOU

